

# Exercise\_1 ver c

July 6, 2019

## 1 Deliverables:

- Submit two files that have been named in the following manner: YourLastName\_Exercise\_1 and that includes the following files:
  - Your PDF document that includes your source code and output.
  - Your ipynb script that includes your source code and output.

## 2 Objectives:

In this exercise, you will:

- Use Jupyter notebook to run ipynb script
- Experiment with file processing (reading and writing files in different formats)
- Use different data types to store and process data

Formatting Python Code When programming in Python, refer to Kenneth Reitz' PEP 8: The Style Guide for Python Code: <http://pep8.org/> (Links to an external site.)Links to an external site. There is the Google style guide for Python at <https://google.github.io/styleguide/pyguide.html> (Links to an external site.)Links to an external site. Comment often and in detail. There are many different kinds of data to be managed and analyzed today, and there are many ways to do it using Python. Being able to manage and modify data isn't useful unless you can also get data into Python, and save your results from it. Beginning with this session we're going to review techniques for input and output from Python starting with the simplest file formats and some basic Python tools. We'll also take a first look at the pandas package. Pandas has become very popular amongst Pythonic data scientists and is being used at the largest of the big data firms. In the sessions that follow we'll consider more complicated file types and data munging tools and techniques. So, let's start with flat files. A flat file is just a file that's, well, flat. It's typically a string of characters that may include end of line markers like a newline or carriage return code. Let's write a simple flat file out to disk by entering the following command:

```
[10]: ### https://www.dataquest.io/blog/jupyter-notebook-tips-tricks-shortcuts/
### Execute the code line by line in jupyter-notebook
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
```

```
[5]: outfile = open('myflatfile.txt', 'w') # open to write to a text file
```

```
[5]: <_io.TextIOWrapper name='myflatfile.txt' mode='w' encoding='cp1252'>
```

Note: In [1]: represents the command prompt in your IPython session. Depending on what you've been doing in a session, the digit or digits you see in it will vary. But you knew that, right?

```
[2]: outfile # outfile is an open file 'object' in write mode. By default it's a
      ↪text, not binary, file
```

```
[2]: <_io.TextIOWrapper name='myflatfile.txt' mode='w' encoding='cp1252'>
```

```
[3]: type(outfile)
```

```
[3]: _io.TextIOWrapper
```

Pardon the slight digression, above. It's purpose was to show what kind of Python 'object' outfile is. (Everything in Python is an object, right?) the .txt file name extension is optional. Now, let's create a text string and then write it to outfile:

```
[4]: iLikeButter='''Slather me toast with a bargefull of butter, and crown it with a
      ↪bucket of Pythonberry jam.'''
```

This is obviously a quote from a high cholesterol data science pirate. How many characters are in this string? Try the function len(iLikeButter).

```
[5]: len(iLikeButter)
```

```
[5]: 91
```

Write the string to outfile and then close outfile:

```
[6]: outfile.write(iLikeButter)
```

```
[6]: 91
```

```
[7]: outfile.close() # it's good practice to close whatever you open
```

```
[8]: import os
      os.getcwd()
```

```
[8]: 'C:\\Users\\asidd\\Desktop\\MSDS\\420 Database Systems\\Lecture 2\\Exercise 1
      Files ver c\\Exercise 1 Files ver c'
```

Let's read your file back in:

```
[11]: infile=open('myflatfile.txt','r') #read as a text file. For a binary,
      ↪you'd use 'rb'
      doYouWantButter=infile.read() #reads the file contents into a string
      ↪variable
      doYouWantButter #this should give you the iLikeButter
      ↪string value
      type(doYouWantButter) # this should give you str
```

```
[11]: 'Slather me toast with a bargefull of butter, and crown it with a bucket of
      Pythonberry jam.'
```

```
[11]: str
```

Next, let's read a text file with more than one line. The text file louielouie.txt has been provided to you. Pop it into the default directory for your session, the directory you identified before. Then,

open it for reading:

```
[12]: kingsMenLouie=open('louielouie.txt','r')           #'r' since this is a  
      →text file.
```

```
[13]: louielouie=kingsMenLouie.read()                   # take a look at louielouie by typing  
      →its name
```

And, then you could split the lines in louielouie into a list of lines as strings:

```
[14]: louielist=louielouie.split()                      # lists are your Python Friend.  
      →(One of them, at least)
```

You could also read this file line by line with `readline()`. For example, to get the file contents into louie a string variable (and assuming that the file has been closed and opened again after the foregoing):

```
[15]: louie=""                                           # string var where we're going to put the lines from the file
```

```
[16]: while True:  
      line=kingsMenLouie.readline()  
      if not line:  
          break  
      louie+=line
```

Give the above code a try to see what you get. Are `louie` and `louielouie` different? Try the command `louie==louielouie`.

Python usually does a good job closing files that have been opened, but it's good practice to do so explicitly whenever possible. This is especially true when you are writing data out to a file, as explicitly closing a file written to forces any remaining write operation to finish. Did you close all the files you opened, above?

A simple way to close a file you've written to is as follows. Suppose you want to write the character string `iLikeButter` to a file called `greaseitup.txt` in your current directory. If you do:

```
[19]: louie==louielouie
```

```
[19]: False
```

```
[15]: with open('greaseitup.txt','wt') as butterOut:  
      butterOut.write(iLikeButter)
```

the file will be closed automatically for you when your write operation is completed. Note the `'wt'` in the open statement. `'t'` is for text, but it's optional. if you include a `'b'` instead, you'll have a binary file instead of a text file.

The procedures for reading and writing binary files using `open`, `.read`, etc. are for the most part the same as for text files, and so we're not going to spend time here on binary file input and output. We're shortly going to move on to reading and writing csv files, but before that let's take a look at the classic method for serializing (storing with permanency) python objects called pickling. A pickle file includes one or more Python objects that can be read back into Python that has a Python-specific, environment independent format.

```
[20]: import pickle
```

Now let's pickle our `louielist` from above in a file in the current working directory.

```
[21]: pickle.dump(louielist,open('louielist.p','wb'))
```

The above writes a binary pickle file. You can read the file back into Python like:

```
[22]: louielsBack=pickle.load(open('louielist.p','rb'))
```

Did you get louielist back unchanged? Try `louielist == louielsBack` from the command prompt.

We're going to move ahead to consider csv files, but to do so we're going to make use of the pandas package. So let's import pandas first, and then look at a simple example of a very useful panda object, the DataFrame.

```
[23]: import pandas as pd          # panda's nickname is pd

import numpy as np              # numpy as np

from pandas import DataFrame, Series    # for convenience
```

My guess is that you have used the numpy package before in your work or in a previous course. DataFrame and Series are very handy pandas data structures that can do yeoman work for you in your data management efforts.

By way of introduction, let's first read a little pickled pandas DataFrame. A pandas DataFrame is a table-like data structure with columns that can be of different data types, and that has both row and column indices.

A Series is like one column of a DataFrame. It's a kind of vector that has an associated index. A DataFrame can be thought of as a set of Series in the columns that share a single index, the row index.

DataFrames and Series have many useful attributes and features, some of which we'll explore in upcoming exercises. But now let's try reading a less trivial csv file into a DataFrame. The file is xyzcust10.csv, and it should be available to you on Canvas. Take a look at it with your favorite text editor. Then, put it in a place you can find it from Canvas, and input it into a DataFrame:

```
[24]: xyzcust10=pd.read_csv('xyzcust10.csv')
```

The file has 10 variables in it. The rows, or records, are XYZ customers. How many records are in xyzcust10?

What types of variables are in the columns of xyzcust10? To find out:

```
[26]: xyzcust10.dtypes
```

```
[26]: ACCTNO          object
      ZIP            int64
      ZIP4           int64
      LTD_SALES      float64
      LTD_TRANSACTIONS  int64
      YTD_SALES_2009  float64
      YTD_TRANSACTIONS_2009 int64
      CHANNEL_ACQUISITION object
      BUYER_STATUS   object
      ZIP9_Supercode  int64
      ZIP9_SUPERCODE  int64
      dtype: object
```

Data Dictionary for xycust10

image.png

Look at the first and last rows in xyzcust10:

[70]: xyzcust10.head()

```
[70]:      ACCTNO      ZIP  ZIP4  LTD_SALES  LTD_TRANSACTIONS  YTD_SALES_2009  \
0  WDQQLLDQL  60084  5016      90.0           1           0.0
1  WQWAYHYLA  60091  1750     4227.0           9          1263.0
2  GSHAPLHAW  60067   900      420.0           3           129.0
3  PGGYDYWAD  60068  3838     6552.0           6           0.0
4  LWPSGPLLS  60090  3932      189.0           3           72.0

      YTD_TRANSACTIONS_2009  CHANNEL_ACQUISITION  BUYER_STATUS  ZIP9_Supercode  \
0                        0                    IB      INACTIVE    600845016
1                        3                    RT      ACTIVE     600911750
2                        1                    RT      ACTIVE     600670900
3                        0                    RT      INACTIVE    600683838
4                        1                    RT      ACTIVE     600903932

      ZIP9_SUPERCODE
0      600845016
1      600911750
2      600670900
3      600683838
4      600903932
```

[28]: xyzcust10.tail()

```
[28]:      ACCTNO      ZIP  ZIP4  LTD_SALES  LTD_TRANSACTIONS  YTD_SALES_2009  \
30466  SYDQYLSWH  60098  3951     2736.0           9           96.0
30467  SAPDQHQLP  60098  9681     2412.0           8          108.0
30468  SASYPDSY  60098   0       429.0           1           0.0
30469  PWQPDWHA  60098  7927      651.0           1           0.0
30470  SQQHDIYHWH  60098  4160     4527.0          16          672.0

      YTD_TRANSACTIONS_2009  CHANNEL_ACQUISITION  BUYER_STATUS  ZIP9_Supercode  \
30466                        1                    RT      ACTIVE    600983951
30467                        1                    RT      ACTIVE    600989681
30468                        0                    RT      INACTIVE    600983858
30469                        0                    RT      INACTIVE    600987927
30470                        2                    RT      ACTIVE    600984160

      ZIP9_SUPERCODE
30466      600983951
30467      600989681
30468      600983858
30469      600987927
30470      600984160
```

Note that in this file missing values for ZIP, ZIP4, and the nine digit ZIP are represented with zeros, 0's. The ZIPs could really also be coded as strings, rather than as integers, couldn't they? Also, it looks like there might be two nine digit ZIP code variables. Are they the same?

That is are the values in these two variables the same for every row of data? How would you locate the rows in xyzcust10 that have a zero for ZIP or for ZIP4? We'll see in the next session's Python Practice.

### 3 Requirements :

1. Produce a list of Zip values in xyzcust10 along with their frequencies
2. How many records with missing ZIP in xyzcust10?
3. How many active and inactive BUYER\_STATUS in xyzcust10?
4. Of all the active customers, what is the break down based upon the channel acquisition?
5. Check that the data matches the data dictionary. Are all Lapsed customers properly labeled - meaning they did not purchase anything in 2009? Show your code to how you came to your answer.
6. List the top 10 'inactive' customers by life to date sales?

```
[81]: # Write your python code here with one new cell for each requirement
# 1 Produce a list of Zip values in xyzcust10 along with their frequencies
xyzcust10.ZIP.value_counts()
```

```
[81]: 60091      3458
      60093      3178
      60062      3099
      60067      3050
      60068      2781
      60089      2007
      60056      1529
      60074      1313
      60060      1296
      60061      1207
      60076      1090
      60069       784
      60077       740
      60084       723
      60073       686
      60090       648
      60098       564
      60070       463
      60085       379
      60083       344
      60081       322
      60087       268
      60097       151
      60096       125
      60071        98
```

```

60064      42
60072      34
60088      28
60078      25
60065      21
60075       5
60094       4
60082       3
60079       2
60192       2
60095       1
0          1
Name: ZIP, dtype: int64

```

```

[82]: #2 How many records with missing ZIP in xyzcust10?
print(sum(xyzcust10.ZIP==0), "records with missing ZIP in xyzcust10")
print(xyzcust10.loc[:, 'ZIP'].isnull().sum(), "null(s) in Zip column")

```

```

1 records with missing ZIP in xyzcust10
0 null(s) in Zip column

```

```

[108]: #3 How many active and inactive BUYER_STATUS in xyzcust10?
xyzcust10.BUYER_STATUS.value_counts()

print('There are {0} active and {1} inactive BUYER_STATUS in xyzcust10.'.format(
    xyzcust10.BUYER_STATUS[xyzcust10.BUYER_STATUS=='ACTIVE'].count(),
    xyzcust10.BUYER_STATUS[xyzcust10.BUYER_STATUS=='INACTIVE'].count()))

```

```

[108]: ACTIVE      13465
      INACTIVE     9078
      LAPSED       7928
Name: BUYER_STATUS, dtype: int64

```

There are 13465 active and 9078 inactive BUYER\_STATUS in xyzcust10

```

[84]: #4 Of all the active customers, what is the break down based upon the channel_
      ↳ acquisition?
xyzcust10.CHANNEL_ACQUISITION[xyzcust10.BUYER_STATUS=='ACTIVE'].value_counts()

```

```

[84]: RT      10827
      IB      1829
      CB       809
Name: CHANNEL_ACQUISITION, dtype: int64

```

```

[109]: #5 Check that the data matches the data dictionary. Are all Lapsed customers_
      ↳ properly labeled -
      # meaning they did not purchase anything in 2009? Show your code to how you_
      ↳ came to your answer.

```

```
## sum(xyzcust10.BUYER_STATUS=='LAPSED')
sum(xyzcust10.YTD_TRANSACTIONS_2009[xyzcust10.BUYER_STATUS=='LAPSED'])
#or
sum(xyzcust10.YTD_SALES_2009[xyzcust10.BUYER_STATUS=='LAPSED'])
print('Yes, all Lapsed customers properly labeled because Lapsed customers did not
      ↪not purchase anything in 2009')
```

[109]: 0

[109]: 0.0

Yes, all Lapsed customers properly labeled because Lapsed customers did not purchase anything in 2009

[110]: *#6 List the top 10 'inactive' customers by life to date sales?*

```
print('Here is the list of the top 10 inactive customers by LTD sales')
xyzcust10.LTD_SALES[xyzcust10.BUYER_STATUS=='INACTIVE'].nlargest(10)
```

Here is the list of the top 10 inactive customers by LTD sales

[110]:

11714	30696.0
5761	14946.0
15700	12078.0
28839	11529.0
13790	10455.0
11095	10095.0
9375	8889.0
5692	8745.0
12919	8661.0
6535	8598.0

Name: LTD\_SALES, dtype: float64