

Siddikov _ Exercise 4 version b

August 16, 2019

1 Deliverables:

- Submit two files that has the name: YourLastName_Exercise_4:
- 1. Your **PDF document** that has your Source code and output
- 2. Your **ipynb script** that has your Source code and output
- 3. You may zip these 2 files and submit

2 Objectives:

In this exercise, you will:

- Analyze the dataset in the given CSV file
- Clean the given dataset
- Load the dataset into sqlite database engine
- Execute different SQL queries

Formatting Python Code When programming in Python, refer to Kenneth Reitz' PEP 8: The Style Guide for Python Code: <http://pep8.org/> (Links to an external site.)Links to an external site. There is the Google style guide for Python at <https://google.github.io/styleguide/pyguide.html> (Links to an external site.)Links to an external site. Comment often and in detail.

2.0.1 Data Preparation

As a data scientist for BestDeal retailer, you have been tasked with improving their revenue and the effectiveness of the marketing campaign of their electronic products. The given dataset has 10,000 records for the purchases of their customers and is used to predict customers shopping patterns and to provide answers for ad-hoc queries. The dataset DirtyData4BestDeal10000.csv is drawn from its database of customers.

```
[1]: import pandas as pd # panda's nickname is pd

import numpy as np # numpy as np

from pandas import DataFrame, Series # for convenience

import sqlalchemy
```

```

from sqlalchemy import create_engine

from sqlalchemy import inspect

%matplotlib inline
# ignore all future warnings
from warnings import simplefilter
simplefilter(action='ignore', category=FutureWarning)
import warnings
warnings.filterwarnings("ignore")

```

2.0.2 Lets read the dirtydata4bestdeal CSV and load into a dataframe object

```
[2]: dirtydata4bestdeal=pd.read_csv('DirtyData4BestDeal10000.csv')
```

```
[3]: # Do you see NaN values below?
```

```
dirtydata4bestdeal.head()
```

```
[3]:
```

	ZipCode	CustomerAge	SamsungTV46LED	SonyTV42LED	XBOX360	DellLaptop	\
0	30134.0	35.0	1	1	1	0	
1	62791.0	43.0	0	1	0	0	
2	60611.0	23.0	1	NaN	0	1	
3	60616.0	56.0	0	1	1	1	
4	30303.0	25.0	1	NaN	0	NaN	

	BoseSoundSystem	BoseHeadSet	SonyHeadSet	iPod	...	\
0	0	1.0	1.0	0.0	...	
1	1	0.0	1.0	0.0	...	
2	0	NaN	1.0	1.0	...	
3	0	0.0	1.0	1.0	...	
4	1	1.0	1.0	0.0	...	

	GalaxyTablet	SurfaceTablet	HPLaptop	HDMICable	SpeakerCable	\
0	1	0.0	1.0	1.0	1.0	
1	1	0.0	1.0	0.0	1.0	
2	0	0.0	1.0	0.0	1.0	
3	0	0.0	1.0	0.0	1.0	
4	1	0.0	1.0	1.0	1.0	

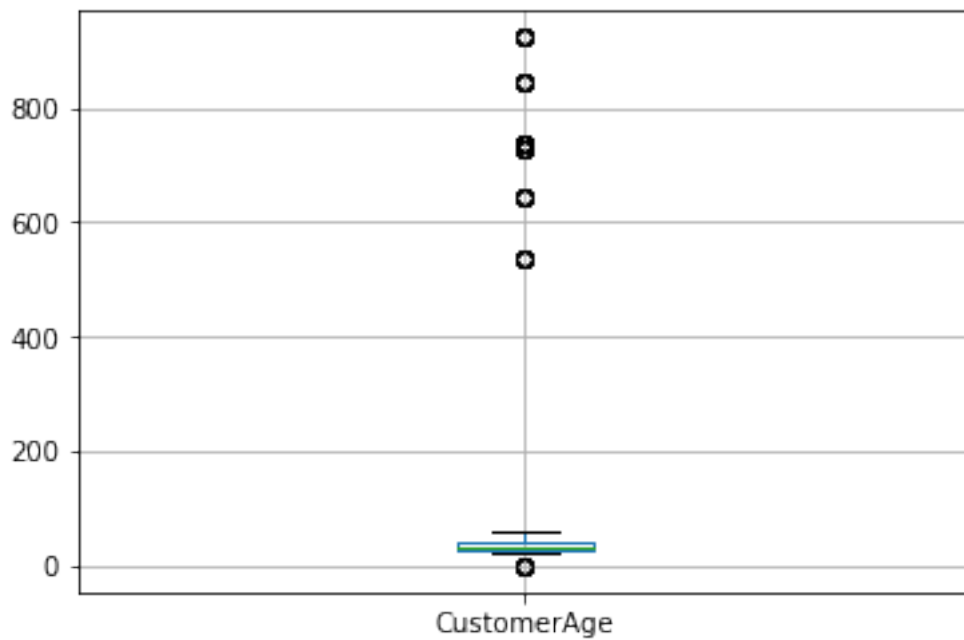
	CallOfDutyGame	GrandTheftAutoGame	ASUSLaptop	LenovoLaptop	\
0	1.0	0.0	1.0	1.0	
1	1.0	0.0	1.0	1.0	
2	1.0	0.0	NaN	1.0	
3	0.0	0.0	1.0	0.0	
4	1.0	0.0	1.0	10.0	

	TVStandWallMount
0	1
1	1
2	1
3	0
4	0


[5 rows x 34 columns]

2.0.3 Lets use boxplot to visualize the data and get an idea if there are dirty/messy/invalid data

```
[4]: # check out customer age
dirtydata4bestdeal.boxplot(column='CustomerAge');
```




```
[5]: # check out customer age with a histogram
dirtydata4bestdeal['CustomerAge'].hist(bins=20);
```



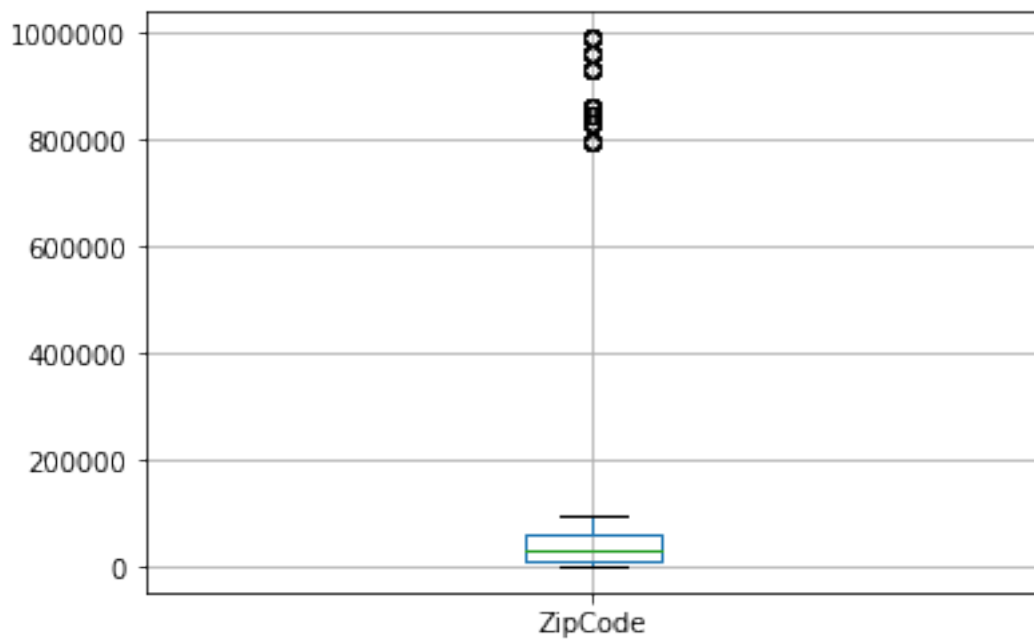
```
[6]: # look at details of LenevoLaptop  
dirtydata4bestdeal.LenevoLaptop.describe()
```

```
[6]: count    9976.000000  
     mean      0.629711  
     std      0.627375  
     min      0.000000  
     25%      0.000000  
     50%      1.000000  
     75%      1.000000  
     max      11.000000  
     Name: LenevoLaptop, dtype: float64
```

```
[7]: dirtydata4bestdeal.boxplot(column='LenevoLaptop');
```



```
[8]: # look at zip codes  
dirtydata4bestdeal.boxplot(column='ZipCode');
```



2.0.4 Lets clean the dirty/messy data in the dirtydata4bestdeal dataframe object

You need to write your python code such that: 1. rows/records/tuples/transactions in the data frame that have missing values for fields/columns will be removed 2. rows/records/tuples/transactions in the data frame that have invalid/abnormal values for fields/columns will be removed

Examples of invalid/dirty/messy data: 1. NaN values in the dataframe (Blank/Empty cells in the CSV file)

2. Every product has a value 1 which means bought or 0 which means NOT bought; values like 11, 10, 9 are examples of invalid data

3. CustomerAge value range could be from 18 to 150; values like 723, 634 are examples of invalid data

4. Zipcode should have at least 5 digits

```
[9]: dirtydata4bestdeal.head()
```

```
[9]:   ZipCode  CustomerAge  SamsungTV46LED  SonyTV42LED  XBOX360  DellLaptop  \
0  30134.0          35.0              1              1          1            0
1  62791.0          43.0              0              1          0            0
2  60611.0          23.0              1             NaN          0            1
3  60616.0          56.0              0              1          1            1
4  30303.0          25.0              1             NaN          0           NaN

   BoseSoundSystem  BoseHeadSet  SonyHeadSet  iPod  ...  \
0                0           1.0          1.0  0.0  ...
1                1           0.0          1.0  0.0  ...
2                0           NaN          1.0  1.0  ...
3                0           0.0          1.0  1.0  ...
4                1           1.0          1.0  0.0  ...

   GalaxyTablet  SurfaceTablet  HPLaptop  HDMICable  SpeakerCable  \
0              1             0.0        1.0        1.0          1.0
1              1             0.0        1.0        0.0          1.0
2              0             0.0        1.0        0.0          1.0
3              0             0.0        1.0        0.0          1.0
4              1             0.0        1.0        1.0          1.0

   CallOfDutyGame  GrandTheftAutoGame  ASUSLaptop  LenevoLaptop  \
0              1.0                  0.0          1.0          1.0
1              1.0                  0.0          1.0          1.0
2              1.0                  0.0          NaN          1.0
3              0.0                  0.0          1.0          0.0
4              1.0                  0.0          1.0         10.0

   TVStandWallMount
0                  1
1                  1
2                  1
3                  0
```

4

0

[5 rows x 34 columns]

[10]: `dirtydata4bestdeal.shape`

[10]: (10000, 34)

[11]: `#`

```

→ -----

# Add the rest of your code here to clean the data

# steps you must take
# - eliminate NA's
# - product values should only be either a 0 or a 1
# - customer's age needs to be valid
# - zipcodes should have at least 5 digits

# Optional steps
# - if there are other things you want to clean, clearly document them
#   and run them in this section before you create a database

#
→ -----

```

[12]: `# Drop the NaN values`

```

cleandata4bestdeal=dirtydata4bestdeal.dropna()
cleandata4bestdeal.head()

# Do you see NaN values dropped below?

```

```

[12]:   ZipCode  CustomerAge  SamsungTV46LED  SonyTV42LED  XBOX360  DellLaptop  \
0  30134.0         35.0             1             1           1             0
1  62791.0         43.0             0             1           0             0
3  60616.0         56.0             0             1           1             1
5   2108.0         55.0             1             1           1             1
6  90033.0         44.0             1             1           1             1

      BoseSoundSystem  BoseHeadSet  SonyHeadSet  iPod  ...  \
0                   0           1.0           1.0  0.0  ...
1                   1           0.0           1.0  0.0  ...
3                   0           0.0           1.0  1.0  ...
5                  10           0.0           0.0  0.0  ...
6                   0           0.0           0.0  0.0  ...

      GalaxyTablet  SurfaceTablet  HPLaptop  HDMICable  SpeakerCable  \

```

0	1	0.0	1.0	1.0	1.0
1	1	0.0	1.0	0.0	1.0
3	0	0.0	1.0	0.0	1.0
5	1	1.0	1.0	1.0	1.0
6	1	1.0	1.0	1.0	0.0

	CallOfDutyGame	GrandTheftAutoGame	ASUSLaptop	LenevoLaptop	\
0	1.0	0.0	1.0	1.0	
1	1.0	0.0	1.0	1.0	
3	0.0	0.0	1.0	0.0	
5	1.0	0.0	1.0	0.0	
6	1.0	1.0	0.0	0.0	

	TVStandWallMount
0	1
1	1
3	0
5	0
6	1

[5 rows x 34 columns]

```
[13]: cleandata4bestdeal.shape
```

```
[13]: (9432, 34)
```

```
[14]: # convert objects and floats into integers
cleandata4bestdeal['SonyTV42LED'] = pd.
    ↳to_numeric(cleandata4bestdeal['SonyTV42LED'], errors='coerce').fillna(0).
    ↳astype(int)
cleandata4bestdeal['XBOX360'] = pd.to_numeric(cleandata4bestdeal['XBOX360'],
    ↳errors='coerce').fillna(0).astype(int)
cleandata4bestdeal['DellLaptop'] = pd.
    ↳to_numeric(cleandata4bestdeal['DellLaptop'], errors='coerce').fillna(0).
    ↳astype(int)
cleandata4bestdeal['BoseSoundSystem'] = pd.
    ↳to_numeric(cleandata4bestdeal['BoseSoundSystem'], errors='coerce').fillna(0).
    ↳astype(int)

cleandata4bestdeal = cleandata4bestdeal.astype('int32')
```

```
[15]: # product values should only be either a 0 or a 1
##cleandata4bestdeal.loc[:, 'SamsungTV46LED': 'TVStandWallMount'] = \
##cleandata4bestdeal.loc[:, 'SamsungTV46LED': 'TVStandWallMount']\
##[cleandata4bestdeal.loc[:, 'SamsungTV46LED': 'TVStandWallMount']\
##.isin([0, 1])].fillna(0).astype(int)

cleandata4bestdeal = \
cleandata4bestdeal[(cleandata4bestdeal.iloc[:, 2:] <= 1).all(1)]
```



```
[16]: cleandata4bestdeal.shape
```

```
[16]: (9206, 34)
```

```
[17]: cleandata4bestdeal.loc[:, 'SamsungTV46LED': 'TVStandWallMount'].apply(pd.  
      ↪value_counts)
```

```
[17]: SamsungTV46LED  SonyTV42LED  XBOX360  DellLaptop  BoseSoundSystem  \  
0           3067           1762           1742           4436           4763  
1           6139           7444           7464           4770           4443  
  
BoseHeadSet  SonyHeadSet  iPod  iPhone  Panasonic50LED  ...  \  
0           4501           1490  7620    5924           7018  ...  
1           4705           7716  1586    3282           2188  ...  
  
GalaxyTablet  SurfaceTablet  HPLaptop  HDMICable  SpeakerCable  \  
0           2814           8567         NaN         4717         2984  
1           6392           639    9206.0         4489         6222  
  
CallOfDutyGame  GrandTheftAutoGame  ASUSLaptop  LenevoLaptop  \  
0           2537                   6066         3821         3518  
1           6669                   3140         5385         5688  
  
TVStandWallMount  
0           2659  
1           6547  
  
[2 rows x 32 columns]
```

```
[18]: #There are ages zero and over 500; which are invalid  
cleandata4bestdeal.loc[:, 'CustomerAge'].value_counts().sort_index()
```

```
[18]: 0           8  
21        201  
22        267  
23       735  
24       184  
25       472  
26       343  
27       505  
28       494  
29       462  
30        16  
31       210  
32       184  
33       168  
34       480  
35       373  
36       192
```

37	511
38	457
39	104
41	96
42	183
43	464
44	528
45	128
46	151
47	104
49	184
51	81
53	24
54	296
55	66
56	184
57	144
59	119
61	32
536	8
643	8
727	8
737	16
843	8
923	8

Name: CustomerAge, dtype: int64

[19]: *#customer's age needs to be valid*

```
cleandata4bestdeal_1 = cleandata4bestdeal[
    cleandata4bestdeal['CustomerAge'].between(20,100)]

cleandata4bestdeal_1.loc[:, 'CustomerAge'].value_counts().sort_index()
```

[19]:

21	201
22	267
23	735
24	184
25	472
26	343
27	505
28	494
29	462
30	16
31	210
32	184
33	168
34	480

35	373
36	192
37	511
38	457
39	104
41	96
42	183
43	464
44	528
45	128
46	151
47	104
49	184
51	81
53	24
54	296
55	66
56	184
57	144
59	119
61	32

Name: CustomerAge, dtype: int64

[20]: *# zipcodes should have at least 5 digits*

```
cleandata4bestdeal.loc[:, 'ZipCode'].astype('int').value_counts().sort_index()
```

[20]:

2108	613
2109	930
2110	224
10065	762
30134	1141
30303	985
33129	539
33130	280
44114	510
60532	243
60585	240
60603	224
60611	62
60616	960
62791	3
90024	144
90033	639
94102	164
94158	495
794158	8
830134	8

```
844114      8
860616      8
960616      8
990033      8
Name: ZipCode, dtype: int64
```

```
[21]: #zipcodes should have at least 5 digits
```

```
cleandata4bestdeal_2 = cleandata4bestdeal_1[
    cleandata4bestdeal_1['ZipCode'] < 100000]
cleandata4bestdeal_2['ZipCode'].value_counts().sort_index()
```

```
[21]: 2108      613
      2109      918
      2110      224
      10065     750
      30134    1133
      30303     985
      33129     531
      33130     280
      44114     510
      60532     243
      60585     240
      60603     224
      60611      62
      60616     952
      62791      3
      90024     144
      90033     631
      94102     164
      94158     487
Name: ZipCode, dtype: int64
```


```
[22]: # check the df shape after cleaning the data
```

```
print(cleandata4bestdeal.shape)
print(cleandata4bestdeal_2.shape)
```


```
(9206, 34)
(9094, 34)
```

```
[ ]:
```

```
[23]: # after cleaning the customer age - does the boxplot still show outliers?
      # how does the histogram look?
      # if this does not look better - you are not ready to proceed
      cleandata4bestdeal_2.boxplot(column='CustomerAge');
```




```
[24]: cleandata4bestdeal_2['CustomerAge'].hist(bins=20);
```



```
[25]: # boxplot after cleaning the zip code
```

```
cleandata4bestdeal_2.boxplot(column='ZipCode');
```



2.0.5 Lets store the cleaned data into the Database

```
[26]: # how many records did you end up with after the data cleaning?  
cleandata4bestdeal = cleandata4bestdeal_2  
cleandata4bestdeal.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 9094 entries, 0 to 9999  
Data columns (total 34 columns):  
ZipCode          9094 non-null int32  
CustomerAge      9094 non-null int32  
SamsungTV46LED   9094 non-null int32  
SonyTV42LED      9094 non-null int32  
XBOX360          9094 non-null int32  
DellLaptop       9094 non-null int32  
BoseSoundSystem  9094 non-null int32  
BoseHeadSet      9094 non-null int32  
SonyHeadSet      9094 non-null int32  
iPod             9094 non-null int32  
iPhone           9094 non-null int32  
Panasonic50LED   9094 non-null int32  
SonyPS4          9094 non-null int32  
WiiU             9094 non-null int32
```

```

WDexternalHD          9094 non-null int32
SamsungTV55LED        9094 non-null int32
SonyTV60LED           9094 non-null int32
SandiskMemoryCard     9094 non-null int32
SonySoundSystem       9094 non-null int32
SonyCamera            9094 non-null int32
PanasonicCamera       9094 non-null int32
HPPrinter             9094 non-null int32
SonyDVDplayer         9094 non-null int32
ToshibaDVDplayer      9094 non-null int32
GalaxyTablet          9094 non-null int32
SurfaceTablet         9094 non-null int32
HPLaptop              9094 non-null int32
HDMICable             9094 non-null int32
SpeakerCable          9094 non-null int32
CallOfDutyGame        9094 non-null int32
GrandTheftAutoGame    9094 non-null int32
ASUSLaptop            9094 non-null int32
LenevoLaptop          9094 non-null int32
TVStandWallMount      9094 non-null int32
dtypes: int32(34)
memory usage: 1.2 MB

```

```

[27]: # now that your data has been cleaned, lets store it in a database

# NOTE - if you run this code more than once, the database will exist and this
→section will fail
# NOTE - to run this more than once, you need to delete the database first
#       OR - change the database name to create a new database

engine = create_engine('sqlite:///bestdeal1.db')

```

```

[28]: cleandata4bestdeal.to_sql('trans4cust', engine)

```

```

    ** Sanity Test: Did it create the table in bestdeal.db? Check!!**

```

```

[29]: insp=inspect(engine)

```

```

[30]: insp.get_table_names()

```

```

[30]: ['trans4cust']

```

```

[31]: pd.read_sql_table('trans4cust', engine).columns

```

```

[31]: Index(['index', 'ZipCode', 'CustomerAge', 'SamsungTV46LED', 'SonyTV42LED',
'XBOX360', 'DellLaptop', 'BoseSoundSystem', 'BoseHeadSet',
'SonyHeadSet', 'iPod', 'iPhone', 'Panasonic50LED', 'SonyPS4', 'WiiU',
'WDexternalHD', 'SamsungTV55LED', 'SonyTV60LED', 'SandiskMemoryCard',
'SonySoundSystem', 'SonyCamera', 'PanasonicCamera', 'HPPrinter',
'SonyDVDplayer', 'ToshibaDVDplayer', 'GalaxyTablet', 'SurfaceTablet',
'HPLaptop', 'HDMICable', 'SpeakerCable', 'CallOfDutyGame',

```

```
'GrandTheftAutoGame', 'ASUSLaptop', 'LenevoLaptop', 'TVStandWallMount'],
dtype='object')
```

should produce the columns of the DataFrame you wrote to the db.

2.0.6 Now we are ready to query the Database

Query example #1: get the transactions for the customers in zipCode 60616

```
[32]: # =====
# *****
#
# WARNING - this pre-run notebook is using dirty data
# WARNING - after cleaning the data, your output should look different
#
# =====
# *****
```

```
[33]: resultsForBestDealCustTrans=pd.read_sql_query("SELECT * FROM trans4cust WHERE_
↳ZipCode='60616'", engine)
```

```
[34]: resultsForBestDealCustTrans.head()
```

```
[34]:   index  ZipCode  CustomerAge  SamsungTV46LED  SonyTV42LED  XBOX360  \
0      3    60616           56              0             1          1
1     16    60616           43              0             1          1
2     18    60616           54              1             0          0
3     23    60616           43              1             1          1
4     34    60616           31              0             1          1

      DellLaptop  BoseSoundSystem  BoseHeadSet  SonyHeadSet  ...  \
0              1                0            0            1    ...
1              0                1            0            1    ...
2              1                0            1            1    ...
3              0                1            1            1    ...
4              1                0            0            1    ...

      GalaxyTablet  SurfaceTablet  HPLaptop  HDMICable  SpeakerCable  \
0                0                0          1          0            1
1                1                0          1          1            1
2                0                1          1          0            1
3                1                1          1          1            0
4                1                0          1          1            1

      CallOfDutyGame  GrandTheftAutoGame  ASUSLaptop  LenevoLaptop  \
0                  0                    0          1            0
1                  1                    0          1            1
2                  1                    0          1            1
3                  1                    0          1            1
4                  1                    1          0            0
```


	TVStandWallMount
0	0
1	1
2	1
3	1
4	1

[5 rows x 35 columns]

Query example #2: get the transactions for ALL customers

```
[35]: resultsForBestDealCustTrans=pd.read_sql_query("SELECT * \
FROM trans4cust", engine)
```

```
[36]: resultsForBestDealCustTrans.head()
```

```
[36]:   index  ZipCode  CustomerAge  SamsungTV46LED  SonyTV42LED  XBOX360  \
0      0    30134         35             1             1           1
1      1    62791         43             0             1           0
2      3    60616         56             0             1           1
3      6    90033         44             1             1           1
4      9     2109         37             0             1           1
```

	DellLaptop	BoseSoundSystem	BoseHeadSet	SonyHeadSet	...	\
0	0	0	1	1	...	
1	0	1	0	1	...	
2	1	0	0	1	...	
3	1	0	0	0	...	
4	0	1	0	1	...	

	GalaxyTablet	SurfaceTablet	HPLaptop	HDMICable	SpeakerCable	\
0	1	0	1	1	1	
1	1	0	1	0	1	
2	0	0	1	0	1	
3	1	1	1	1	0	
4	0	0	1	0	1	

	CallOfDutyGame	GrandTheftAutoGame	ASUSLaptop	LenovoLaptop	\
0	1	0	1	1	
1	1	0	1	1	
2	0	0	1	0	
3	1	1	0	0	
4	0	0	1	1	

	TVStandWallMount
0	1
1	1

2	0
3	1
4	0

[5 rows x 35 columns]

Query example #3: get the number of customers in every ZipCode sorted by ZipCode

```
[37]: resultsForBestDealCustTrans=pd.read_sql_query("SELECT ZipCode , COUNT(*) as num_
        ↳customers' \
            FROM trans4cust \
            GROUP BY ZipCode \
            ORDER BY ZipCode", engine)
```

```
[38]: resultsForBestDealCustTrans
```

```
[38]:   ZipCode  num_customers
0      2108             613
1      2109             918
2      2110             224
3     10065             750
4     30134            1133
5     30303             985
6     33129             531
7     33130             280
8     44114             510
9     60532             243
10    60585             240
11    60603             224
12    60611              62
13    60616            952
14    62791              3
15    90024            144
16    90033            631
17    94102            164
18    94158            487
```

Query example #4: get the number of customers for every Age Group in ZipCode 60616 sorted by CustomerAge

```
[39]: resultsForBestDealCustTrans=pd.read_sql_query(
        "SELECT CustomerAge , COUNT(*) as 'num_customers' \
        FROM trans4cust \
        WHERE ZipCode=60616 \
        GROUP BY CustomerAge \
        ORDER BY CustomerAge", engine)
```

```
[40]: resultsForBestDealCustTrans
```

```
[40]: CustomerAge  num_customers
0      21          56
1      22          32
2      23          40
3      25          88
4      26          48
5      27          32
6      28          32
7      29          56
8      31          16
9      32          16
10     34          96
11     35          72
12     37          64
13     38          24
14     39           8
15     43          48
16     44          88
17     45          24
18     46          24
19     51           8
20     54          48
21     56          32
```

Query example #5: Plot in a stacked-bar figure the number of customers who bought SonyTV60LED and/or BoseSoundSystem in every zipcode that has more than 400 customers who bought these two products(either bought one of these products or the two products)

```
[41]: SonyTV60LEDCustTrans=pd.read_sql_query(
      "SELECT ZipCode , COUNT(*) as 'num_customers' FROM trans4cust \
      WHERE SonyTV60LED=1 GROUP BY ZipCode HAVING COUNT(*) > 400", engine)

      BoseSoundSystemCustTrans=pd.read_sql_query(
      "SELECT ZipCode , COUNT(*) as 'num_customers' FROM trans4cust \
      WHERE BoseSoundSystem=1 GROUP BY ZipCode HAVING COUNT(*) > 400", engine)
```

```
[42]: SonyTV60LEDCustTrans
```

```
[42]: ZipCode  num_customers
0      2108          402
1      2109          579
2      10065          439
3      30134          757
4      30303          517
5      60616          689
```

```
[43]: BoseSoundSystemCustTrans
```

```
[43]: ZipCode  num_customers
0      2109           424
1      30134          799
2      30303          464
3      60616          466
4      90033          404
```

```
[44]: SonyTV60LEDCustTrans.ZipCode
```

```
[44]: 0      2108
1      2109
2      10065
3      30134
4      30303
5      60616
Name: ZipCode, dtype: int64
```

```
[45]: import numpy

#   There are zipcodes that Sony got bought but not Bose
#   but there are also zipcodes that Bose got bought but not Sony
#
#   AND we need to use stacked-bar graph and we have a potentially asymmetrical
#   →set of zipcode values
#   So, we need to do somework to create the symmteric set of zipcode values
#   →for Sony and Bose

sonyZipCodeTuples=tuple(SonyTV60LEDCustTrans.ZipCode.astype(numpy.int))
sony_num_customersTuples=tuple(SonyTV60LEDCustTrans.num_customers.astype(numpy.
    →int))

boseZipCodeTuples=tuple(BoseSoundSystemCustTrans.ZipCode.astype(numpy.int))
bose_num_customersTuples=tuple(BoseSoundSystemCustTrans.num_customers.
    →astype(numpy.int))

sony_dict = dict(zip(sonyZipCodeTuples, sony_num_customersTuples))
bose_dict = dict(zip(boseZipCodeTuples, bose_num_customersTuples))

for key in bose_dict.keys():
    if ((key in sony_dict.keys()) == False): sony_dict[key]=0

for key in sony_dict.keys():
    if ((key in bose_dict.keys()) == False): bose_dict[key]=0

bose_zip= sorted(bose_dict.keys())

sony_zip= sorted(sony_dict.keys())
```

```

bose_zip_tuple=tuple(bose_zip)

sony_zip_tuple=tuple(sony_zip)

bose_customer_list=[]

for bose in bose_zip_tuple:
    bose_customer_list.append(bose_dict[bose])

sony_customer_list=[]

for sony in sony_zip_tuple:
    sony_customer_list.append(sony_dict[sony])

bose_customer_tuple=tuple(bose_customer_list)
sony_customer_tuple=tuple(sony_customer_list)

```

```

[46]: # See docs for bar_stack at the URL
      # http://matplotlib.org/examples/pylab_examples/bar_stacked.html

import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline

ind = np.arange(len(sony_customer_tuple))

# the width of the bars: can also be len(x) sequence
width = .5

p1 = plt.bar(ind, sony_customer_tuple, width, color='r')
p2 = plt.bar(ind, bose_customer_tuple, width, color='y',
             ↳bottom=sony_customer_tuple)

plt.ylabel('Number of Customers')
plt.xlabel('Zip Code')


plt.title('Number of Customers by ZipCode and 2 Products')

plt.xticks(ind + width, sony_zip_tuple, horizontalalignment='right')

plt.yticks(np.arange(0, 2000, 100))
plt.legend((p1[0], p2[0]), ('Sony', 'Bose'))

plt.show()

```



3 Requirements :

1. (Use SQL/SQLite): show the top 3 zip codes with the most customers
2. (Use SQL/SQLite): selecting the customers from the top 3 zip codes (results from question 1), what are ages of the customers? Sort output by most customers. You can show all 3 zip codes combined or show ages by zip codes.
3. (Use SQL/SQLite): get the number of customers who bought DellLaptop and HPPrinter for every Age group sorted by CustomerAge.
4. (Use SQL/SQLite): Get the list of ZipCodes where no customer bought XBOX360 (this query means NOT even a single customer in that zip code bought XBOX360).
5. (Use SQL/SQLite/Matplotlib): Plot in a stacked-bar figure the number of customers who bought HPLaptop and/or HPPrinter but did NOT buy WDexternalHD for every Customer-Age group that has more than 100 customers who bought these two products (either bought one of these products or the two products but didn't buy WDexternalHD).

[47]: *# Write your python code that meets the above requirements in this cell*
Question 1
(Use SQL/SQLite): show the top 3 zip codes with the most customers

```
Ans_1 = pd.read_sql_query(
    "SELECT ZipCode, COUNT(*) as 'num_customers' \
    FROM trans4cust \
    GROUP BY ZipCode \
```

```
ORDER BY num_customers DESC\
Limit 3;", engine)
```

Ans_1

```
[47]:   ZipCode  num_customers
0      30134             1133
1      30303             985
2      60616             952
```

```
[48]: # Question 2
# (Use SQL/SQLite): selecting the customers from the top 3 zip codes
# (results from question 1), what are ages of the customers?
# Sort output by most customers. You can show all 3 zip codes combined
# or show ages by zip codes.
```

```
Ans_2 = pd.read_sql_query(
    "SELECT ZipCode, CustomerAge, COUNT(*) as 'num_customers' \
FROM trans4cust \
WHERE ZipCode in (30134, 30303, 60616)\
GROUP BY ZipCode, CustomerAge \
ORDER BY num_customers DESC", engine)
```

Ans_2

```
[48]:   ZipCode  CustomerAge  num_customers
0      30134            25             154
1      60616            34              96
2      60616            25              88
3      60616            44              88
4      30134            29              84
5      30303            26              83
6      30303            27              81
7      30303            44              77
8      30134            43              75
9      30134            34              74
10     30303            29              74
11     30303            23              73
12     60616            35              72
13     30303            43              68
14     30134            28              67
15     60616            37              64
16     30303            34              61
17     30134            32              58
18     30134            44              58
19     30134            37              56
20     60616            21              56
21     60616            29              56
22     30303            41              49
23     30134            22              48
```

24	30134	31	48
25	60616	26	48
26	60616	43	48
27	60616	54	48
28	30134	38	45
29	30303	38	42
..
47	30303	37	28
48	30303	54	28
49	30303	24	27
50	30303	31	27
51	30303	49	27
52	30303	61	27
53	30303	56	26
54	60616	38	24
55	60616	45	24
56	60616	46	24
57	30303	59	21
58	30134	45	19
59	60616	31	16
60	60616	32	16
61	30303	33	14
62	30303	45	14
63	30303	55	14
64	30134	42	10
65	30134	51	9
66	30134	54	8
67	60616	39	8
68	60616	51	8
69	30303	30	7
70	30303	35	7
71	30303	51	7
72	30134	26	2
73	30134	33	2
74	30134	39	2
75	30134	59	2
76	30134	46	1

[77 rows x 3 columns]

```
[49]: # Question 3
# (Use SQL/Sqlite): get the number of customers who bought DellLaptop
# and HPPrinter for every Age group sorted by CustomerAge.

Ans_3 = pd.read_sql_query(
"SELECT CustomerAge, COUNT(*) as 'num_customers' \
FROM trans4cust \
```



```
WHERE DellLaptop = 1\
AND HPPrinter = 1\
GROUP BY CustomerAge \
ORDER BY CustomerAge", engine)
```

Ans_3

```
[49]: CustomerAge  num_customers
0          21          201
1          22          203
2          23          304
3          25           64
4          26          183
5          27          272
6          28           56
7          29          143
8          31          194
9          32          184
10         34          120
11         35          136
12         36          192
13         38           16
14         39           88
15         42           72
16         44          184
17         45           32
18         46           63
19         47           32
20         51           16
21         53           24
22         54          127
23         56          176
24         57           64
25         59           80
26         61           32
```

```
[50]: # Question 4
# (Use SQL/SQLite): Get the list of ZipCodes where no customer bought XBOX360
# (this query means NOT even a single customer in that zip code bought XBOX360).
```

```
Ans_4 = pd.read_sql_query(
"SELECT ZipCode, COUNT(*) as 'num_customers' \
FROM trans4cust \
WHERE XBOX360 = 0\
GROUP BY ZipCode", engine)
```

Ans_4

```
[50]: ZipCode  num_customers
0      2108           49
```

1	2109	210
2	2110	96
3	10065	164
4	30134	248
5	30303	220
6	33129	67
7	33130	40
8	44114	81
9	60532	32
10	60585	96
11	60603	88
12	60611	8
13	60616	81
14	62791	3
15	90024	16
16	90033	87
17	94102	36
18	94158	112

```
[51]: # Question 5
# (Use SQL/Sqlite/Matplotlib): Plot in a stacked-bar figure the number of
# →customers
# who bought HPLaptop and/or HPPrinter but did NOT buy WDEternalHD for every
# →CustomerAge
# group that has more than 100 customers who bought these two products
# (either bought one of these products or the two products but didn't buy
# →WDEternalHD).
```

```
Ans_5 = pd.read_sql_query(
"SELECT CustomerAge , COUNT(*) as 'num_customers' \
FROM trans4cust \
WHERE HPLaptop = 1 \
AND HPPrinter = 1 \
AND WDEternalHD = 0 \
GROUP BY CustomerAge HAVING COUNT(*) > 100", engine)
```

```
Ans_5_HPLaptop = pd.read_sql_query(
"SELECT CustomerAge , COUNT(*) as 'num_customers' \
FROM trans4cust \
WHERE HPLaptop = 1 \
AND WDEternalHD = 0 \
GROUP BY CustomerAge HAVING COUNT(*) > 100", engine)
```

```
Ans_5_HPPrinter = pd.read_sql_query(
"SELECT CustomerAge , COUNT(*) as 'num_customers' \
FROM trans4cust \
WHERE HPPrinter = 1 \
```

```
AND WDEternalHD = 0 \
GROUP BY CustomerAge HAVING COUNT(*) > 100", engine)
```

```
[52]: Ans_5_HPLaptop
```

```
[52]:      CustomerAge  num_customers
0           21           192
1           22           203
2           23           437
3           25           147
4           26           205
5           27           307
6           28           126
7           29           290
8           31           196
9           32           176
10          34           178
11          35           348
12          36           183
13          42           177
14          43           104
15          44           319
16          54           148
17          56           162
```

```
[53]: Ans_5_HPPrinter
```

```
[53]:      CustomerAge  num_customers
0           21           192
1           22           193
2           23           437
3           25           147
4           26           205
5           27           288
6           28           126
7           29           290
8           31           196
9           32           176
10          34           156
11          35           348
12          36           183
13          42           177
14          44           292
15          54           121
16          56           162
```

```
[54]: HPLaptop_cusage = tuple(Ans_5_HPLaptop.CustomerAge.astype(numpy.int))
      HPLaptop_cus = tuple(Ans_5_HPLaptop.num_customers.astype(numpy.int))
```

```

HPPrinter_cusage = tuple(Ans_5_HPPrinter.CustomerAge.astype(numpy.int))
HPPrinter_cus = tuple(Ans_5_HPPrinter.num_customers.astype(numpy.int))

HPLaptop_dict = dict(zip(HPLaptop_cusage, HPLaptop_cus))
HPPrinter_dict = dict(zip(HPPrinter_cusage, HPPrinter_cus))

for key in HPPrinter_dict.keys():
    if ((key in HPLaptop_dict.keys()) == False): HPLaptop_dict[key]=0

for key in HPLaptop_dict.keys():
    if ((key in HPPrinter_dict.keys()) == False): HPPrinter_dict[key]=0

HPPrinter_cusage= sorted(HPPrinter_dict.keys())

HPLaptop_cusage= sorted(HPLaptop_dict.keys())

HPPrinter_cusage_tuple=tuple(HPPrinter_cusage)

HPLaptop_cusage_tuple=tuple(HPLaptop_cusage)

HPPrinter_customer_list=[]

for HPPrinter in HPPrinter_cusage_tuple:
    HPPrinter_customer_list.append(HPPrinter_dict[HPPrinter])

HPLaptop_customer_list=[]

for HPLaptop in HPLaptop_cusage_tuple:
    HPLaptop_customer_list.append(HPLaptop_dict[HPLaptop])

HPPrinter_customer_tuple=tuple(HPPrinter_customer_list)
HPLaptop_customer_tuple=tuple(HPLaptop_customer_list)

```

```

[55]: import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline

ind = np.arange(len(HPLaptop_customer_tuple))

# the width of the bars: can also be len(x) sequence
width = 0.5

p1 = plt.bar(ind, HPLaptop_customer_tuple, width, color='r')
p2 = plt.bar(ind, HPPrinter_customer_tuple, width, color='y', bottom =_
→HPLaptop_customer_tuple)

```

```

plt.ylabel('Number of Purchases')
plt.xlabel('Customer Age')

plt.title('Number of Customers by Customer Age and 2 Products')

plt.xticks(ind + width, HPLaptop_cusage_tuple, horizontalalignment='right')

plt.yticks(np.arange(0, 1000, 100))
plt.legend((p1[0], p2[0]), ('HPLaptop', 'HPPrinter'))

plt.show();

```

