

Siddikov _ Exercise 4 version b

August 16, 2019

1 Deliverables:

- Submit two files that has the name: YourLastName_Exercise_4:
- 1. Your **PDF document** that has your Source code and output
- 2. Your **ipynb script** that has your Source code and output
- 3. You may zip these 2 files and submit

2 Objectives:

In this exercise, you will:

- Analyze the dataset in the given CSV file
- Clean the given dataset
- Load the dataset into sqlite database engine
- Execute different SQL queries

Formatting Python Code When programming in Python, refer to Kenneth Reitz' PEP 8: The Style Guide for Python Code: <http://pep8.org/> (Links to an external site.)Links to an external site. There is the Google style guide for Python at <https://google.github.io/styleguide/pyguide.html> (Links to an external site.)Links to an external site. Comment often and in detail.

2.0.1 Data Preparation

As a data scientist for BestDeal retailer, you have been tasked with improving their revenue and the effectiveness of the marketing campaign of their electronic products. The given dataset has 10,000 records for the purchases of their customers and is used to predict customers shopping patterns and to provide answers for ad-hoc queries. The dataset DirtyData4BestDeal10000.csv is drawn from its database of customers.

```
[1]: import pandas as pd # panda's nickname is pd

import numpy as np # numpy as np

from pandas import DataFrame, Series # for convenience

import sqlalchemy
```

```

from sqlalchemy import create_engine

from sqlalchemy import inspect

%matplotlib inline
# ignore all future warnings
from warnings import simplefilter
simplefilter(action='ignore', category=FutureWarning)
import warnings
warnings.filterwarnings("ignore")

```

2.0.2 Lets read the dirtydata4bestdeal CSV and load into a dataframe object

```
[2]: dirtydata4bestdeal=pd.read_csv('DirtyData4BestDeal10000.csv')
```

```
[3]: # Do you see NaN values below?
```

```
dirtydata4bestdeal.head()
```

```
[3]:
```

| | ZipCode | CustomerAge | SamsungTV46LED | SonyTV42LED | XBOX360 | DellLaptop | \ |
|---|---------|-------------|----------------|-------------|---------|------------|---|
| 0 | 30134.0 | 35.0 | 1 | 1 | 1 | 0 | |
| 1 | 62791.0 | 43.0 | 0 | 1 | 0 | 0 | |
| 2 | 60611.0 | 23.0 | 1 | NaN | 0 | 1 | |
| 3 | 60616.0 | 56.0 | 0 | 1 | 1 | 1 | |
| 4 | 30303.0 | 25.0 | 1 | NaN | 0 | NaN | |

| | BoseSoundSystem | BoseHeadSet | SonyHeadSet | iPod | ... | \ |
|---|-----------------|-------------|-------------|------|-----|---|
| 0 | 0 | 1.0 | 1.0 | 0.0 | ... | |
| 1 | 1 | 0.0 | 1.0 | 0.0 | ... | |
| 2 | 0 | NaN | 1.0 | 1.0 | ... | |
| 3 | 0 | 0.0 | 1.0 | 1.0 | ... | |
| 4 | 1 | 1.0 | 1.0 | 0.0 | ... | |

| | GalaxyTablet | SurfaceTablet | HPLaptop | HDMICable | SpeakerCable | \ |
|---|--------------|---------------|----------|-----------|--------------|---|
| 0 | 1 | 0.0 | 1.0 | 1.0 | 1.0 | |
| 1 | 1 | 0.0 | 1.0 | 0.0 | 1.0 | |
| 2 | 0 | 0.0 | 1.0 | 0.0 | 1.0 | |
| 3 | 0 | 0.0 | 1.0 | 0.0 | 1.0 | |
| 4 | 1 | 0.0 | 1.0 | 1.0 | 1.0 | |

| | CallOfDutyGame | GrandTheftAutoGame | ASUSLaptop | LenovoLaptop | \ |
|---|----------------|--------------------|------------|--------------|---|
| 0 | 1.0 | 0.0 | 1.0 | 1.0 | |
| 1 | 1.0 | 0.0 | 1.0 | 1.0 | |
| 2 | 1.0 | 0.0 | NaN | 1.0 | |
| 3 | 0.0 | 0.0 | 1.0 | 0.0 | |
| 4 | 1.0 | 0.0 | 1.0 | 10.0 | |


| | TVStandWallMount |
|---|------------------|
| 0 | 1 |
| 1 | 1 |
| 2 | 1 |
| 3 | 0 |
| 4 | 0 |

[5 rows x 34 columns]

2.0.3 Lets use boxplot to visualize the data and get an idea if there are dirty/messy/invalid data


```
[4]: # check out customer age
dirtydata4bestdeal.boxplot(column='CustomerAge')
```

```
[4]: <matplotlib.axes._subplots.AxesSubplot at 0x25111d66a20>
```



```
[5]: # check out customer age with a histogram
dirtydata4bestdeal['CustomerAge'].hist(bins=20)
```

```
[5]: <matplotlib.axes._subplots.AxesSubplot at 0x25111e31320>
```




```
[6]: # look at details of LenevoLaptop  
dirtydata4bestdeal.LenevoLaptop.describe()
```

```
[6]: count    9976.000000  
     mean      0.629711  
     std      0.627375  
     min      0.000000  
     25%      0.000000  
     50%      1.000000  
     75%      1.000000  
     max     11.000000  
     Name: LenevoLaptop, dtype: float64
```


```
[7]: dirtydata4bestdeal.boxplot(column='LenevoLaptop')
```

```
[7]: <matplotlib.axes._subplots.AxesSubplot at 0x251140dcac8>
```



```
[8]: # look at zip codes
dirtydata4bestdeal.boxplot(column='ZipCode')
```

```
[8]: <matplotlib.axes._subplots.AxesSubplot at 0x25111e264a8>
```



2.0.4 Lets clean the dirty/messy data in the dirtydata4bestdeal dataframe object

You need to write your python code such that: 1. rows/records/tuples/transactions in the data frame that have missing values for fields/columns will be removed 2. rows/records/tuples/transactions in the data frame that have invalid/abnormal values for fields/columns will be removed

Examples of invalid/dirty/messy data: 1. NaN values in the dataframe (Blank/Empty cells in the CSV file)

2. Every product has a value 1 which means bought or 0 which means NOT bought; values like 11, 10, 9 are examples of invalid data

3. CustomerAge value range could be from 18 to 150; values like 723, 634 are examples of invalid data

4. Zipcode should have at least 5 digits

```
[9]: dirtydata4bestdeal.head()
```

```
[9]:  ZipCode  CustomerAge  SamsungTV46LED  SonyTV42LED  XBOX360  DellLaptop  \
0  30134.0          35.0              1              1          1          0
1  62791.0          43.0              0              1          0          0
2  60611.0          23.0              1             NaN          0          1
3  60616.0          56.0              0              1          1          1
4  30303.0          25.0              1             NaN          0         NaN

    BoseSoundSystem  BoseHeadSet  SonyHeadSet  iPod  ...  \
0                0            1.0          1.0  0.0  ...
1                1            0.0          1.0  0.0  ...
2                0            NaN          1.0  1.0  ...
3                0            0.0          1.0  1.0  ...
4                1            1.0          1.0  0.0  ...

    GalaxyTablet  SurfaceTablet  HPLaptop  HDMICable  SpeakerCable  \
0              1             0.0          1.0          1.0          1.0
1              1             0.0          1.0          0.0          1.0
2              0             0.0          1.0          0.0          1.0
3              0             0.0          1.0          0.0          1.0
4              1             0.0          1.0          1.0          1.0

    CallOfDutyGame  GrandTheftAutoGame  ASUSLaptop  LenevoLaptop  \
0              1.0                  0.0          1.0          1.0
1              1.0                  0.0          1.0          1.0
2              1.0                  0.0          NaN          1.0
3              0.0                  0.0          1.0          0.0
4              1.0                  0.0          1.0         10.0

    TVStandWallMount
0                  1
1                  1
2                  1
3                  0
```

[5 rows x 34 columns]

```
[10]: #
# -----
# Add the rest of your code here to clean the data

# steps you must take
# - eliminate NA's
# - product values should only be either a 0 or a 1
# - customer's age needs to be valid
# - zipcodes should have at least 5 digits

# Optional steps
# - if there are other things you want to clean, clearly document them
#   and run them in this section before you create a database

#
# -----
```

```
[11]: # Drop the NaN values

cleandata4bestdeal=dirtydata4bestdeal.dropna()
cleandata4bestdeal.head()

# Do you see NaN values dropped below?
```

```
[11]: ZipCode  CustomerAge  SamsungTV46LED  SonyTV42LED  XBOX360  DellLaptop  \
0  30134.0      35.0           1           1           1           0
1  62791.0      43.0           0           1           0           0
3  60616.0      56.0           0           1           1           1
5   2108.0      55.0           1           1           1           1
6  90033.0      44.0           1           1           1           1

BoseSoundSystem  BoseHeadSet  SonyHeadSet  iPod  ...  \
0                0           1.0         1.0  0.0  ...
1                1           0.0         1.0  0.0  ...
3                0           0.0         1.0  1.0  ...
5               10           0.0         0.0  0.0  ...
6                0           0.0         0.0  0.0  ...

GalaxyTablet  SurfaceTablet  HPLaptop  HDMICable  SpeakerCable  \
0              1           0.0         1.0         1.0         1.0
1              1           0.0         1.0         0.0         1.0
3              0           0.0         1.0         0.0         1.0
```

| | | | | | |
|---|---|-----|-----|-----|-----|
| 5 | 1 | 1.0 | 1.0 | 1.0 | 1.0 |
| 6 | 1 | 1.0 | 1.0 | 1.0 | 0.0 |

| | CallOfDutyGame | GrandTheftAutoGame | ASUSLaptop | LenevoLaptop | \ |
|---|----------------|--------------------|------------|--------------|---|
| 0 | 1.0 | 0.0 | 1.0 | 1.0 | |
| 1 | 1.0 | 0.0 | 1.0 | 1.0 | |
| 3 | 0.0 | 0.0 | 1.0 | 0.0 | |
| 5 | 1.0 | 0.0 | 1.0 | 0.0 | |
| 6 | 1.0 | 1.0 | 0.0 | 0.0 | |

| | TVStandWallMount |
|---|------------------|
| 0 | 1 |
| 1 | 1 |
| 3 | 0 |
| 5 | 0 |
| 6 | 1 |

[5 rows x 34 columns]

```
[12]: # convert objects and floats into integers
cleandata4bestdeal['SonyTV42LED'] = pd.
    →to_numeric(cleandata4bestdeal['SonyTV42LED'], errors='coerce').fillna(0).
    →astype(int)
cleandata4bestdeal['XBOX360'] = pd.to_numeric(cleandata4bestdeal['XBOX360'],
    →errors='coerce').fillna(0).astype(int)
cleandata4bestdeal['DellLaptop'] = pd.
    →to_numeric(cleandata4bestdeal['DellLaptop'], errors='coerce').fillna(0).
    →astype(int)
cleandata4bestdeal['BoseSoundSystem'] = pd.
    →to_numeric(cleandata4bestdeal['BoseSoundSystem'], errors='coerce').fillna(0).
    →astype(int)

cleandata4bestdeal = cleandata4bestdeal.astype('int32')
```

```
[13]: # product values should only be either a 0 or a 1
cleandata4bestdeal.loc[:, 'SamsungTV46LED': 'TVStandWallMount'] = \
cleandata4bestdeal.loc[:, 'SamsungTV46LED': 'TVStandWallMount']\
[cleandata4bestdeal.loc[:, 'SamsungTV46LED': 'TVStandWallMount'].isin([0, 1])].
    →fillna(0).astype(int)
```

```
[14]: cleandata4bestdeal.loc[:, 'SamsungTV46LED': 'TVStandWallMount'].apply(pd.
    →value_counts)
```

| | SamsungTV46LED | SonyTV42LED | XBOX360 | DellLaptop | BoseSoundSystem | \ |
|---|----------------|-------------|---------|------------|-----------------|---|
| 0 | 3123 | 1834 | 1814 | 4519 | 4907 | |
| 1 | 6309 | 7598 | 7618 | 4913 | 4525 | |

| | BoseHeadSet | SonyHeadSet | iPod | iPhone | Panasonic50LED | ... | \ |
|---|-------------|-------------|------|--------|----------------|-----|---|
| 0 | 4595 | 1536 | 7806 | 6030 | 7215 | ... | |


```

1          4837          7896 1626    3402          2217          ...

    GalaxyTablet  SurfaceTablet  HPLaptop  HDMICable  SpeakerCable  \
0          2884          8773      NaN      4812          3073
1          6548          659    9432.0      4620          6359

    CallOfDutyGame  GrandTheftAutoGame  ASUSLaptop  LenevoLaptop  \
0          2569          6204      3927          3638
1          6863          3228      5505          5794

    TVStandWallMount
0          2745
1          6687

```

[2 rows x 32 columns]

```

[15]: #There are zero and over 500 ages; which are invalid
      cleandata4bestdeal.loc[:, 'CustomerAge'].value_counts().sort_index()

```

```

[15]: 0          8
      21         209
      22         281
      23         775
      24         184
      25         497
      26         352
      27         513
      28         495
      29         472
      30          16
      31         224
      32         184
      33         184
      34         488
      35         373
      36         200
      37         512
      38         457
      39         104
      41          96
      42         184
      43         472
      44         536
      45         128
      46         151
      47         104
      49         192
      51         104

```

| | |
|-----|-----|
| 53 | 24 |
| 54 | 297 |
| 55 | 72 |
| 56 | 192 |
| 57 | 144 |
| 59 | 120 |
| 61 | 32 |
| 536 | 8 |
| 643 | 8 |
| 727 | 8 |
| 737 | 16 |
| 843 | 8 |
| 923 | 8 |

Name: CustomerAge, dtype: int64

[16]: *#Zip codes should be less than 5 digits*

```
cleandata4bestdeal.loc[:, 'ZipCode'].astype('int').value_counts().sort_index()
```

[16]:

| | |
|--------|------|
| 2108 | 632 |
| 2109 | 955 |
| 2110 | 224 |
| 10065 | 788 |
| 30134 | 1173 |
| 30303 | 1001 |
| 33129 | 554 |
| 33130 | 280 |
| 44114 | 526 |
| 60532 | 243 |
| 60585 | 248 |
| 60603 | 240 |
| 60611 | 62 |
| 60616 | 960 |
| 62791 | 3 |
| 90024 | 144 |
| 90033 | 665 |
| 94102 | 166 |
| 94158 | 512 |
| 794158 | 8 |
| 830134 | 8 |
| 844114 | 8 |
| 860616 | 8 |
| 930134 | 8 |
| 960616 | 8 |
| 990033 | 8 |

Name: ZipCode, dtype: int64

```
[17]: #customer's age needs to be valid
```

```
cleandata4bestdeal_1 = cleandata4bestdeal[  
    cleandata4bestdeal['CustomerAge'].between(20,70)]
```

```
[18]: #zipcodes should have at least 5 digits
```

```
cleandata4bestdeal_2 = cleandata4bestdeal_1[  
    cleandata4bestdeal_1['ZipCode'] < 100000]
```

```
[19]: # check the df shape after cleaning the data
```

```
print(cleandata4bestdeal.shape)  
print(cleandata4bestdeal_2.shape)
```

```
(9432, 34)
```

```
(9312, 34)
```

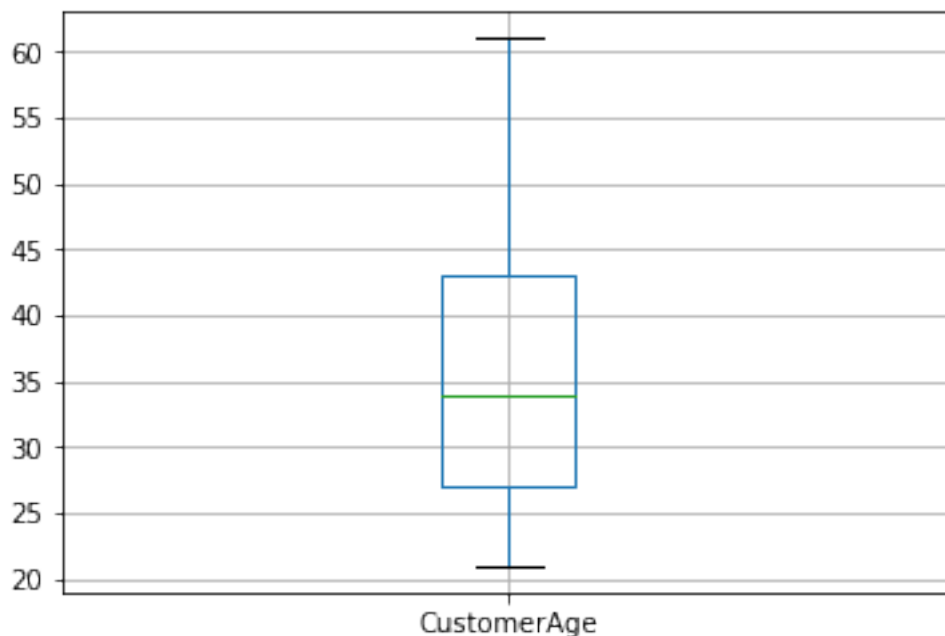
```
[ ]:
```

```
[20]: # after cleaning the customer age - does the boxplot still show outliers?
```


```
# how does the histogram look?
```

```
# if this does not look better - you are not ready to proceed
```


```
cleandata4bestdeal_2.boxplot(column='CustomerAge');
```



```
[21]: cleandata4bestdeal_2['CustomerAge'].hist(bins=20);
```



```
[22]: # boxplot after cleaning the zip code  
cleandata4bestdeal_2.boxplot(column='ZipCode');
```



2.0.5 Lets store the cleaned data into the Database

```
[23]: # how many records did you end up with after the data cleaning?  
cleandata4bestdeal = cleandata4bestdeal_2  
cleandata4bestdeal.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 9312 entries, 0 to 9999  
Data columns (total 34 columns):  
ZipCode          9312 non-null int32  
CustomerAge      9312 non-null int32  
SamsungTV46LED   9312 non-null int32  
SonyTV42LED      9312 non-null int32  
XBOX360          9312 non-null int32  
DellLaptop       9312 non-null int32  
BoseSoundSystem  9312 non-null int32  
BoseHeadSet      9312 non-null int32  
SonyHeadSet      9312 non-null int32  
iPod             9312 non-null int32  
iPhone           9312 non-null int32  
Panasonic50LED   9312 non-null int32  
SonyPS4          9312 non-null int32  
WiiU             9312 non-null int32  
WDexternalHD     9312 non-null int32  
SamsungTV55LED   9312 non-null int32  
SonyTV60LED      9312 non-null int32  
SandiskMemoryCard 9312 non-null int32  
SonySoundSystem  9312 non-null int32  
SonyCamera       9312 non-null int32  
PanasonicCamera  9312 non-null int32  
HPPrinter        9312 non-null int32  
SonyDVDplayer    9312 non-null int32  
ToshibaDVDplayer 9312 non-null int32  
GalaxyTablet     9312 non-null int32  
SurfaceTablet    9312 non-null int32  
HPLaptop         9312 non-null int32  
HDMICable        9312 non-null int32  
SpeakerCable     9312 non-null int32  
CallOfDutyGame   9312 non-null int32  
GrandTheftAutoGame 9312 non-null int32  
ASUSLaptop       9312 non-null int32  
LenevoLaptop     9312 non-null int32  
TVStandWallMount 9312 non-null int32  
dtypes: int32(34)  
memory usage: 1.3 MB
```

```
[24]: # now that your data has been cleaned, lets store it in a database

# NOTE - if you run this code more than once, the database will exist and this
      ↳section will fail
# NOTE - to run this more than once, you need to delete the database first
#       OR - change the database name to create a new database

engine = create_engine('sqlite:///bestdeal1.db')

[25]: cleandata4bestdeal.to_sql('trans4cust', engine)

      ** Sanity Test: Did it create the table in bestdeal.db? Check!!**

[26]: insp=inspect(engine)

[27]: insp.get_table_names()

[27]: ['trans4cust']

[28]: pd.read_sql_table('trans4cust', engine).columns

[28]: Index(['index', 'ZipCode', 'CustomerAge', 'SamsungTV46LED', 'SonyTV42LED',
            'XBOX360', 'DellLaptop', 'BoseSoundSystem', 'BoseHeadSet',
            'SonyHeadSet', 'iPod', 'iPhone', 'Panasonic50LED', 'SonyPS4', 'WiiU',
            'WDexternalHD', 'SamsungTV55LED', 'SonyTV60LED', 'SandiskMemoryCard',
            'SonySoundSystem', 'SonyCamera', 'PanasonicCamera', 'HPPrinter',
            'SonyDVDplayer', 'ToshibaDVDplayer', 'GalaxyTablet', 'SurfaceTablet',
            'HPLaptop', 'HDMICable', 'SpeakerCable', 'CallOfDutyGame',
            'GrandTheftAutoGame', 'ASUSLaptop', 'LenevoLaptop', 'TVStandWallMount'],
            dtype='object')
```

should produce the columns of the DataFrame you wrote to the db.

2.0.6 Now we are ready to query the Database

Query example #1: get the transactions for the customers in zipCode 60616

```
[29]: # =====
# *****
#
# WARNING - this pre-run notebook is using dirty data
# WARNING - after cleaning the data, your output should look different
#
# =====
# *****

[30]: resultsForBestDealCustTrans=pd.read_sql_query("SELECT * FROM trans4cust WHERE
      ↳ZipCode='60616'", engine)

[31]: resultsForBestDealCustTrans.head()

[31]:
```

| | index | ZipCode | CustomerAge | SamsungTV46LED | SonyTV42LED | XBOX360 | \ |
|---|-------|---------|-------------|----------------|-------------|---------|---|
| 0 | 3 | 60616 | 56 | 0 | 1 | 1 | |
| 1 | 16 | 60616 | 43 | 0 | 1 | 1 | |

| | | | | | | |
|---|----|-------|----|---|---|---|
| 2 | 18 | 60616 | 54 | 1 | 0 | 0 |
| 3 | 23 | 60616 | 43 | 1 | 1 | 1 |
| 4 | 34 | 60616 | 31 | 0 | 1 | 1 |

| | | | | | | |
|---|------------|-----------------|-------------|-------------|-----|---|
| | DellLaptop | BoseSoundSystem | BoseHeadSet | SonyHeadSet | ... | \ |
| 0 | 1 | 0 | 0 | 1 | ... | |
| 1 | 0 | 1 | 0 | 1 | ... | |
| 2 | 1 | 0 | 1 | 1 | ... | |
| 3 | 0 | 1 | 1 | 1 | ... | |
| 4 | 1 | 0 | 0 | 1 | ... | |

| | | | | | | |
|---|--------------|---------------|----------|-----------|--------------|---|
| | GalaxyTablet | SurfaceTablet | HPLaptop | HDMICable | SpeakerCable | \ |
| 0 | 0 | 0 | 1 | 0 | 1 | |
| 1 | 1 | 0 | 1 | 1 | 1 | |
| 2 | 0 | 1 | 1 | 0 | 1 | |
| 3 | 1 | 1 | 1 | 1 | 0 | |
| 4 | 1 | 0 | 1 | 1 | 1 | |

| | | | | | |
|---|----------------|--------------------|------------|--------------|---|
| | CallOfDutyGame | GrandTheftAutoGame | ASUSLaptop | LenevoLaptop | \ |
| 0 | 0 | 0 | 1 | 0 | |
| 1 | 1 | 0 | 1 | 1 | |
| 2 | 1 | 0 | 1 | 1 | |
| 3 | 1 | 0 | 1 | 1 | |
| 4 | 1 | 1 | 0 | 0 | |

| | |
|---|------------------|
| | TVStandWallMount |
| 0 | 0 |
| 1 | 1 |
| 2 | 1 |
| 3 | 1 |
| 4 | 1 |

[5 rows x 35 columns]

Query example #2: get the transactions for ALL customers

```
[32]: resultsForBestDealCustTrans=pd.read_sql_query("SELECT * \
FROM trans4cust", engine)
```

```
[33]: resultsForBestDealCustTrans.head()
```

```
[33]:
```

| | index | ZipCode | CustomerAge | SamsungTV46LED | SonyTV42LED | XBOX360 | \ |
|---|-------|---------|-------------|----------------|-------------|---------|---|
| 0 | 0 | 30134 | 35 | 1 | 1 | 1 | |
| 1 | 1 | 62791 | 43 | 0 | 1 | 0 | |
| 2 | 3 | 60616 | 56 | 0 | 1 | 1 | |
| 3 | 5 | 2108 | 55 | 1 | 1 | 1 | |
| 4 | 6 | 90033 | 44 | 1 | 1 | 1 | |

| | DellLaptop | BoseSoundSystem | BoseHeadSet | SonyHeadSet | ... | \ |
|---|------------|-----------------|-------------|-------------|-----|---|
| 0 | 0 | 0 | 1 | 1 | ... | |
| 1 | 0 | 1 | 0 | 1 | ... | |
| 2 | 1 | 0 | 0 | 1 | ... | |
| 3 | 1 | 0 | 0 | 0 | ... | |
| 4 | 1 | 0 | 0 | 0 | ... | |

| | GalaxyTablet | SurfaceTablet | HPLaptop | HDMICable | SpeakerCable | \ |
|---|--------------|---------------|----------|-----------|--------------|---|
| 0 | 1 | 0 | 1 | 1 | 1 | |
| 1 | 1 | 0 | 1 | 0 | 1 | |
| 2 | 0 | 0 | 1 | 0 | 1 | |
| 3 | 1 | 1 | 1 | 1 | 1 | |
| 4 | 1 | 1 | 1 | 1 | 0 | |

| | CallOfDutyGame | GrandTheftAutoGame | ASUSLaptop | LenevoLaptop | \ |
|---|----------------|--------------------|------------|--------------|---|
| 0 | 1 | 0 | 1 | 1 | |
| 1 | 1 | 0 | 1 | 1 | |
| 2 | 0 | 0 | 1 | 0 | |
| 3 | 1 | 0 | 1 | 0 | |
| 4 | 1 | 1 | 0 | 0 | |

| | TVStandWallMount |
|---|------------------|
| 0 | 1 |
| 1 | 1 |
| 2 | 0 |
| 3 | 0 |
| 4 | 1 |

[5 rows x 35 columns]

Query example #3: get the number of customers in every ZipCode sorted by ZipCode

```
[34]: resultsForBestDealCustTrans=pd.read_sql_query("SELECT ZipCode , COUNT(*) as_
    ↳'num_customers' \
        FROM trans4cust \
        GROUP BY ZipCode \
        ORDER BY ZipCode", engine)
```

```
[35]: resultsForBestDealCustTrans
```

```
[35]:   ZipCode  num_customers
0      2108           632
1      2109           943
2      2110           224
3     10065           776
4     30134          1165
5     30303          1001
6     33129           546
```


| | | |
|----|-------|-----|
| 7 | 33130 | 280 |
| 8 | 44114 | 526 |
| 9 | 60532 | 243 |
| 10 | 60585 | 248 |
| 11 | 60603 | 240 |
| 12 | 60611 | 62 |
| 13 | 60616 | 952 |
| 14 | 62791 | 3 |
| 15 | 90024 | 144 |
| 16 | 90033 | 657 |
| 17 | 94102 | 166 |
| 18 | 94158 | 504 |

Query example #4: get the number of customers for every Age Group in ZipCode 60616 sorted by CustomerAge

```
[36]: resultsForBestDealCustTrans=pd.read_sql_query(
      "SELECT CustomerAge , COUNT(*) as 'num_customers' \
        FROM trans4cust \
        WHERE ZipCode=60616 \
        GROUP BY CustomerAge \
        ORDER BY CustomerAge", engine)
```

```
[37]: resultsForBestDealCustTrans
```

```
[37]:
```

| | CustomerAge | num_customers |
|----|-------------|---------------|
| 0 | 21 | 56 |
| 1 | 22 | 32 |
| 2 | 23 | 40 |
| 3 | 25 | 88 |
| 4 | 26 | 48 |
| 5 | 27 | 32 |
| 6 | 28 | 32 |
| 7 | 29 | 56 |
| 8 | 31 | 16 |
| 9 | 32 | 16 |
| 10 | 34 | 96 |
| 11 | 35 | 72 |
| 12 | 37 | 64 |
| 13 | 38 | 24 |
| 14 | 39 | 8 |
| 15 | 43 | 48 |
| 16 | 44 | 88 |
| 17 | 45 | 24 |
| 18 | 46 | 24 |
| 19 | 51 | 8 |
| 20 | 54 | 48 |
| 21 | 56 | 32 |

Query example #5: Plot in a stacked-bar figure the number of customers who bought SonyTV60LED and/or BoseSoundSystem in every zipcode that has more than 400 customers who bought these two products(either bought one of these products or the two products)

```
[38]: SonyTV60LEDCustTrans=pd.read_sql_query(
      "SELECT ZipCode , COUNT(*) as 'num_customers' FROM trans4cust \
        WHERE SonyTV60LED=1 GROUP BY ZipCode HAVING COUNT(*) > 400", engine)

      BoseSoundSystemCustTrans=pd.read_sql_query(
      "SELECT ZipCode , COUNT(*) as 'num_customers' FROM trans4cust \
        WHERE BoseSoundSystem=1 GROUP BY ZipCode HAVING COUNT(*) > 400", engine)
```

```
[39]: SonyTV60LEDCustTrans
```

```
[39]:   ZipCode  num_customers
      0      2108           416
      1      2109           599
      2     10065           455
      3     30134           773
      4     30303           524
      5     60616           689
```

```
[40]: BoseSoundSystemCustTrans
```

```
[40]:   ZipCode  num_customers
      0      2109           428
      1     30134           824
      2     30303           472
      3     60616           466
      4     90033           405
```

```
[41]: SonyTV60LEDCustTrans.ZipCode
```

```
[41]: 0      2108
      1      2109
      2     10065
      3     30134
      4     30303
      5     60616
      Name: ZipCode, dtype: int64
```

```
[42]: import numpy

      #   There are zipcodes that Sony got bought but not Bose
      #   but there are also zipcodes that Bose got bought but not Sony
      #
      #   AND we need to use stacked-bar graph and we have a potentially asymmetrical
      #   →set of zipcode values
      #   So, we need to do somework to create the symmteric set of zipcode values
      #   →for Sony and Bose
```

```

sonyZipCodeTuples=tuple(SonyTV60LEDCustTrans.ZipCode.astype(numpy.int))
sony_num_customersTuples=tuple(SonyTV60LEDCustTrans.num_customers.astype(numpy.
    →int))

boseZipCodeTuples=tuple(BoseSoundSystemCustTrans.ZipCode.astype(numpy.int))
bose_num_customersTuples=tuple(BoseSoundSystemCustTrans.num_customers.
    →astype(numpy.int))

sony_dict = dict(zip(sonyZipCodeTuples, sony_num_customersTuples))
bose_dict = dict(zip(boseZipCodeTuples, bose_num_customersTuples))

for key in bose_dict.keys():
    if ((key in sony_dict.keys()) == False): sony_dict[key]=0

for key in sony_dict.keys():
    if ((key in bose_dict.keys()) == False): bose_dict[key]=0

bose_zip= sorted(bose_dict.keys())

sony_zip= sorted(sony_dict.keys())

bose_zip_tuple=tuple(bose_zip)

sony_zip_tuple=tuple(sony_zip)

bose_customer_list=[]

for bose in bose_zip_tuple:
    bose_customer_list.append(bose_dict[bose])

sony_customer_list=[]

for sony in sony_zip_tuple:
    sony_customer_list.append(sony_dict[sony])

bose_customer_tuple=tuple(bose_customer_list)
sony_customer_tuple=tuple(sony_customer_list)

```

[43]: *# See docs for bar_stack at the URL*
http://matplotlib.org/examples/pylab_examples/bar_stacked.html

```

import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline

```

```

ind = np.arange(len(sony_customer_tuple))

# the width of the bars: can also be len(x) sequence
width = .5

p1 = plt.bar(ind, sony_customer_tuple, width, color='r')
p2 = plt.bar(ind, bose_customer_tuple, width, color='y',
             ↳bottom=sony_customer_tuple)

plt.ylabel('Number of Customers')
plt.xlabel('Zip Code')


plt.title('Number of Customers by ZipCode and 2 Products')

plt.xticks(ind + width, sony_zip_tuple, horizontalalignment='right')

plt.yticks(np.arange(0, 2000, 100))
plt.legend((p1[0], p2[0]), ('Sony', 'Bose'))

plt.show()

```



3 Requirements :

1. (Use SQL/SQLite): show the top 3 zip codes with the most customers
2. (Use SQL/SQLite): selecting the customers from the top 3 zip codes (results from question 1), what are ages of the customers? Sort output by most customers. You can show all 3 zip codes combined or show ages by zip codes.
3. (Use SQL/SQLite): get the number of customers who bought DellLaptop and HPPrinter for every Age group sorted by CustomerAge.
4. (Use SQL/SQLite): Get the list of ZipCodes where no customer bought XBOX360 (this query means NOT even a single customer in that zip code bought XBOX360).
5. (Use SQL/SQLite/Matplotlib): Plot in a stacked-bar figure the number of customers who bought HPLaptop and/or HPPrinter but did NOT buy WDexternalHD for every Customer-Age group that has more than 100 customers who bought these two products(either bought one of these products or the two products but didn't buy WDexternalHD).

```
[44]: # Write your python code that meets the above requirements in this cell
# Question 1
# (Use SQL/SQLite): show the top 3 zip codes with the most customers
```

```
Ans_1 = pd.read_sql_query(
    "SELECT ZipCode, COUNT(*) as 'num_customers' \
    FROM trans4cust \
    GROUP BY ZipCode \
    ORDER BY num_customers DESC\
    Limit 3;", engine)
Ans_1
```

```
[44]:   ZipCode  num_customers
0    30134             1165
1    30303             1001
2    60616              952
```

```
[45]: # Question 2
# (Use SQL/SQLite): selecting the customers from the top 3 zip codes
# (results from question 1), what are ages of the customers?
# Sort output by most customers. You can show all 3 zip codes combined
# or show ages by zip codes.
```

```
Ans_2 = pd.read_sql_query(
    "SELECT ZipCode, CustomerAge, COUNT(*) as 'num_customers' \
    FROM trans4cust \
    WHERE ZipCode in (30134, 30303, 60616)\
    GROUP BY ZipCode, CustomerAge \
    ORDER BY num_customers DESC", engine)
Ans_2
```

```
[45]:   ZipCode  CustomerAge  num_customers
0    30134             25              155
1    60616             34              96
```

| | | | |
|----|-------|-----|-----|
| 2 | 60616 | 25 | 88 |
| 3 | 60616 | 44 | 88 |
| 4 | 30134 | 29 | 84 |
| 5 | 30303 | 26 | 83 |
| 6 | 30303 | 27 | 81 |
| 7 | 30303 | 44 | 77 |
| 8 | 30134 | 43 | 75 |
| 9 | 30303 | 29 | 75 |
| 10 | 30134 | 34 | 74 |
| 11 | 30303 | 23 | 73 |
| 12 | 60616 | 35 | 72 |
| 13 | 30303 | 43 | 68 |
| 14 | 30134 | 28 | 67 |
| 15 | 60616 | 37 | 64 |
| 16 | 30303 | 34 | 61 |
| 17 | 30134 | 32 | 58 |
| 18 | 30134 | 44 | 58 |
| 19 | 30134 | 31 | 56 |
| 20 | 30134 | 37 | 56 |
| 21 | 60616 | 21 | 56 |
| 22 | 60616 | 29 | 56 |
| 23 | 30303 | 41 | 49 |
| 24 | 30134 | 22 | 48 |
| 25 | 60616 | 26 | 48 |
| 26 | 60616 | 43 | 48 |
| 27 | 60616 | 54 | 48 |
| 28 | 30134 | 38 | 45 |
| 29 | 30303 | 25 | 42 |
| .. | ... | ... | ... |
| 47 | 30303 | 37 | 28 |
| 48 | 30303 | 54 | 28 |
| 49 | 30303 | 24 | 27 |
| 50 | 30303 | 31 | 27 |
| 51 | 30303 | 49 | 27 |
| 52 | 30303 | 61 | 27 |
| 53 | 30303 | 56 | 26 |
| 54 | 30134 | 51 | 24 |
| 55 | 60616 | 38 | 24 |
| 56 | 60616 | 45 | 24 |
| 57 | 60616 | 46 | 24 |
| 58 | 30303 | 33 | 21 |
| 59 | 30303 | 59 | 21 |
| 60 | 30134 | 45 | 19 |
| 61 | 60616 | 31 | 16 |
| 62 | 60616 | 32 | 16 |
| 63 | 30303 | 45 | 14 |
| 64 | 30303 | 55 | 14 |

| | | | |
|----|-------|----|----|
| 65 | 30134 | 42 | 10 |
| 66 | 30134 | 54 | 8 |
| 67 | 60616 | 39 | 8 |
| 68 | 60616 | 51 | 8 |
| 69 | 30303 | 30 | 7 |
| 70 | 30303 | 35 | 7 |
| 71 | 30303 | 51 | 7 |
| 72 | 30134 | 26 | 2 |
| 73 | 30134 | 33 | 2 |
| 74 | 30134 | 39 | 2 |
| 75 | 30134 | 59 | 2 |
| 76 | 30134 | 46 | 1 |

[77 rows x 3 columns]

[46]: *# Question 3*
(Use SQL/SQLite): get the number of customers who bought DellLaptop
and HPPrinter for every Age group sorted by CustomerAge.

```
Ans_3 = pd.read_sql_query(
"SELECT CustomerAge, COUNT(*) as 'num_customers' \
  FROM trans4cust \
 WHERE DellLaptop = 1\
 AND HPPrinter = 1\
 GROUP BY CustomerAge \
 ORDER BY CustomerAge", engine)
Ans_3
```

[46]:

| | CustomerAge | num_customers |
|----|-------------|---------------|
| 0 | 21 | 201 |
| 1 | 22 | 217 |
| 2 | 23 | 320 |
| 3 | 25 | 65 |
| 4 | 26 | 192 |
| 5 | 27 | 280 |
| 6 | 28 | 56 |
| 7 | 29 | 151 |
| 8 | 31 | 208 |
| 9 | 32 | 184 |
| 10 | 34 | 128 |
| 11 | 35 | 136 |
| 12 | 36 | 200 |
| 13 | 38 | 16 |
| 14 | 39 | 88 |
| 15 | 42 | 72 |
| 16 | 44 | 192 |
| 17 | 45 | 32 |
| 18 | 46 | 63 |

| | | |
|----|----|-----|
| 19 | 47 | 32 |
| 20 | 51 | 24 |
| 21 | 53 | 24 |
| 22 | 54 | 128 |
| 23 | 56 | 184 |
| 24 | 57 | 64 |
| 25 | 59 | 80 |
| 26 | 61 | 32 |

[47]: *# Question 4*
(Use SQL/Sqlite): Get the list of ZipCodes where no customer bought XBOX360
(this query means NOT even a single customer in that zip code bought XBOX360).

```
Ans_4 = pd.read_sql_query(
"SELECT ZipCode, COUNT(*) as 'num_customers' \
  FROM trans4cust \
 WHERE XBOX360 = 0\
 GROUP BY ZipCode", engine)
Ans_4
```

[47]:

| | ZipCode | num_customers |
|----|---------|---------------|
| 0 | 2108 | 56 |
| 1 | 2109 | 216 |
| 2 | 2110 | 96 |
| 3 | 10065 | 168 |
| 4 | 30134 | 248 |
| 5 | 30303 | 220 |
| 6 | 33129 | 73 |
| 7 | 33130 | 40 |
| 8 | 44114 | 97 |
| 9 | 60532 | 32 |
| 10 | 60585 | 96 |
| 11 | 60603 | 88 |
| 12 | 60611 | 8 |
| 13 | 60616 | 81 |
| 14 | 62791 | 3 |
| 15 | 90024 | 16 |
| 16 | 90033 | 104 |
| 17 | 94102 | 36 |
| 18 | 94158 | 128 |

[48]: *# Question 5*
(Use SQL/Sqlite/Matplotlib): Plot in a stacked-bar figure the number of
→customers
who bought HPLaptop and/or HPPrinter but did NOT buy WDexternalHD for every
→CustomerAge
group that has more than 100 customers who bought these two products


```
# (either bought one of these products or the two products but didn't buy
→WDexternalHD).
```

```
Ans_5 = pd.read_sql_query(
"SELECT CustomerAge , COUNT(*) as 'num_customers' \
  FROM trans4cust \
  WHERE HPLaptop = 1 \
  AND HPPrinter = 1 \
  AND WDexternalHD = 0 \
  GROUP BY CustomerAge HAVING COUNT(*) > 100", engine)
```

```
Ans_5_HPLaptop = pd.read_sql_query(
"SELECT CustomerAge , COUNT(*) as 'num_customers' \
  FROM trans4cust \
  WHERE HPLaptop = 1 \
  AND WDexternalHD = 0 \
  GROUP BY CustomerAge HAVING COUNT(*) > 100", engine)
```

```
Ans_5_HPPrinter = pd.read_sql_query(
"SELECT CustomerAge , COUNT(*) as 'num_customers' \
  FROM trans4cust \
  WHERE HPPrinter = 1 \
  AND WDexternalHD = 0 \
  GROUP BY CustomerAge HAVING COUNT(*) > 100", engine)
```

[49]: Ans_5_HPLaptop

```
[49]:
```

| | CustomerAge | num_customers |
|----|-------------|---------------|
| 0 | 21 | 200 |
| 1 | 22 | 216 |
| 2 | 23 | 462 |
| 3 | 25 | 151 |
| 4 | 26 | 214 |
| 5 | 27 | 314 |
| 6 | 28 | 126 |
| 7 | 29 | 299 |
| 8 | 31 | 210 |
| 9 | 32 | 176 |
| 10 | 34 | 186 |
| 11 | 35 | 348 |
| 12 | 36 | 191 |
| 13 | 42 | 178 |
| 14 | 43 | 104 |
| 15 | 44 | 326 |
| 16 | 54 | 149 |
| 17 | 56 | 169 |

[50]: Ans_5_HPPrinter

```
[50]:      CustomerAge  num_customers
      0           21           200
      1           22           206
      2           23           454
      3           25           151
      4           26           214
      5           27           295
      6           28           126
      7           29           298
      8           31           210
      9           32           176
     10           34           164
     11           35           348
     12           36           191
     13           42           178
     14           44           299
     15           54           122
     16           56           169
```

```
[51]: HPLaptop_cusage = tuple(Ans_5_HPLaptop.CustomerAge.astype(numpy.int))
      HPLaptop_cus = tuple(Ans_5_HPLaptop.num_customers.astype(numpy.int))

      HPPrinter_cusage = tuple(Ans_5_HPPrinter.CustomerAge.astype(numpy.int))
      HPPrinter_cus = tuple(Ans_5_HPPrinter.num_customers.astype(numpy.int))

      HPLaptop_dict = dict(zip(HPLaptop_cusage, HPLaptop_cus))
      HPPrinter_dict = dict(zip(HPPrinter_cusage, HPPrinter_cus))

      for key in HPPrinter_dict.keys():
          if ((key in HPLaptop_dict.keys()) == False): HPLaptop_dict[key]=0

      for key in HPLaptop_dict.keys():
          if ((key in HPPrinter_dict.keys()) == False): HPPrinter_dict[key]=0

      HPPrinter_cusage= sorted(HPPrinter_dict.keys())

      HPLaptop_cusage= sorted(HPLaptop_dict.keys())

      HPPrinter_cusage_tuple=tuple(HPPrinter_cusage)

      HPLaptop_cusage_tuple=tuple(HPLaptop_cusage)

      HPPrinter_customer_list=[]

      for HPPrinter in HPPrinter_cusage_tuple:
          HPPrinter_customer_list.append(HPPrinter_dict[HPPrinter])
```

```

HPLaptop_customer_list=[]

for HPLaptop in HPLaptop_cusage_tuple:
    HPLaptop_customer_list.append(HPLaptop_dict[HPLaptop])

HPPrinter_customer_tuple=tuple(HPPrinter_customer_list)
HPLaptop_customer_tuple=tuple(HPLaptop_customer_list)

```

```

[53]: import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline

ind = np.arange(len(HPLaptop_customer_tuple))

# the width of the bars: can also be len(x) sequence
width = .5

p1 = plt.bar(ind, HPLaptop_customer_tuple, width, color='r')
p2 = plt.bar(ind, HPPrinter_customer_tuple, width, color='y', bottom =  

    ↳HPLaptop_customer_tuple)

plt.ylabel('Number of Customers')
plt.xlabel('Customer Age')

plt.title('Number of Customers by Customer Age and 2 Products')

plt.xticks(ind + width, HPLaptop_cusage_tuple, horizontalalignment='right')

plt.yticks(np.arange(0, 2000, 100))
plt.legend((p1[0], p2[0]), ('HPLaptop', 'HPPrinter'))

plt.show();

```

