

Siddikov Exercise 5 version b

August 16, 2019

Deliverables:

- Submit a single zip-compressed file that has the name: YourLastName_Exercise_5 that has the following files:
 1. Your **PDF document** that has your Source code and output
 2. Your **ipynb script** that has your Source code and output
 3. You can zip these 2 files if you like; use the same naming convention for the zip file.

1 Objectives:

In this exercise, you will:

- Construct hierarchical indexes
- Select and group data to create pivot-tables

Formatting Python Code When programming in Python, refer to Kenneth Reitz' PEP 8: The Style Guide for Python Code: <http://pep8.org/> (Links to an external site.)Links to an external site. There is the Google style guide for Python at <https://google.github.io/styleguide/pyguide.html> (Links to an external site.)Links to an external site. Comment often and in detail.

2 Specifications and Requirements

We're going to use the XYZ data again to construct hierarchical indexes and select, modify, group, and reshape data in a wide variety of ways. The data we want here, which we'll call xyzcustnew, are as follows:

```
[1]: import pandas as pd    # panda's nickname is pd
import numpy as np        # numpy as np
from pandas import DataFrame, Series, Categorical
from sqlalchemy import create_engine

import matplotlib

%matplotlib inline

[2]: engine=create_engine('sqlite:///xyz.db')           # the db is in my current_
      ↪working directory
```

```
[3]: # .info gives same feedback as .dtype and .count
xyzcustnew=pd.read_sql_table('xyzcust',engine)
xyzcustnew.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30179 entries, 0 to 30178
Data columns (total 11 columns):
index                30179 non-null int64
ACCTNO               30179 non-null object
ZIP                  30179 non-null int64
ZIP4                 30179 non-null int64
LTD_SALES            30179 non-null float64
LTD_TRANSACTIONS     30179 non-null int64
YTD_SALES_2009       30179 non-null float64
YTD_TRANSACTIONS_2009 30179 non-null int64
CHANNEL_ACQUISITION  30179 non-null object
BUYER_STATUS         30179 non-null object
ZIP9_SUPERCODE       30179 non-null int64
dtypes: float64(2), int64(6), object(3)
memory usage: 2.5+ MB
```

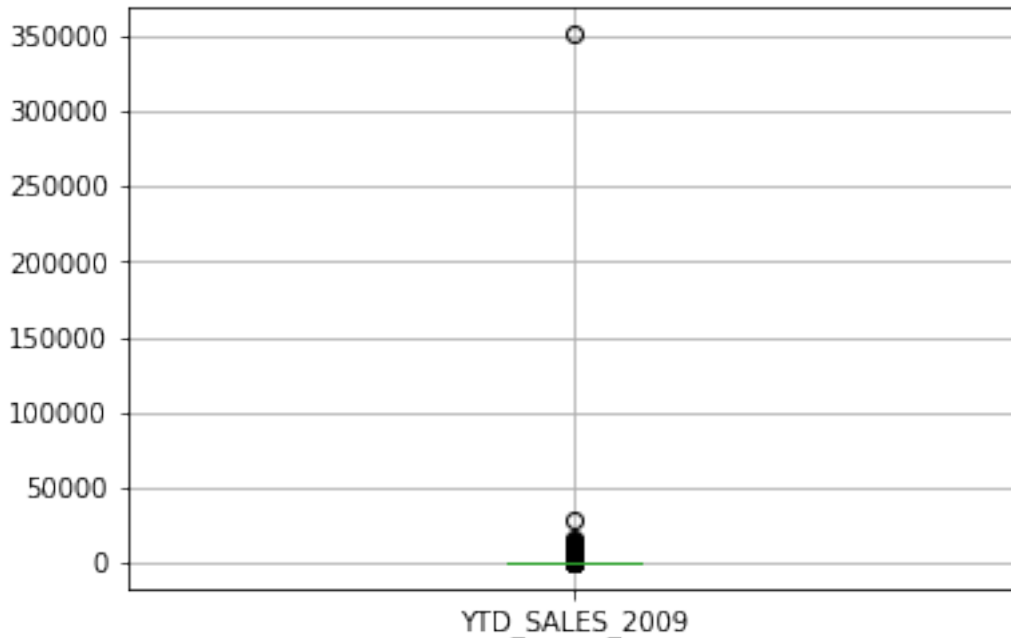
```
[4]: # heavyCut is a constant and was decided as where the data
# should be cut
heavyCut= 423
```

```
[5]: # look at characteristics
xyzcustnew['YTD_SALES_2009'].describe()
```

```
[5]: count      30179.000000
mean         236.283972
std          2117.042293
min           0.000000
25%           0.000000
50%           0.000000
75%          207.000000
max          351000.000000
Name: YTD_SALES_2009, dtype: float64
```

```
[6]: # look at spread of 2009 sales
xyzcustnew.boxplot(column='YTD_SALES_2009')
```

```
[6]: <matplotlib.axes._subplots.AxesSubplot at 0x1fecbfa3b38>
```



```
[7]: # create a categorial variable of either a 1 or 0 based upon the value of YTD_SALES_2009
      # YTD_SALES_2009 greater than the heavyCut value will be assigned a 1
      heavyCat=Categorical(np.where(xyzcustnew.YTD_SALES_2009>heavyCut,1,0))
      heavyCat.describe()
```

```
[7]:      counts      freqs
categories
0          25795  0.854733
1           4384  0.145267
```

```
[8]: # be more descriptive than a 1 or a 0
      heavyCat.rename_categories(['regular','heavy'],inplace=True)
      heavyCat.describe()
```

```
[8]:      counts      freqs
categories
regular      25795  0.854733
heavy         4384  0.145267
```

```
[9]: # look at the first ten records
      heavyCat[:10]
```

```
[9]: [regular, heavy, regular, regular, regular, regular, heavy, regular, regular,
      regular]
      Categories (2, object): [regular, heavy]
```

```
[10]: # create a new column with this variable
      xyzcustnew['heavyCat']=heavyCat
```

```
[11]: # a dummy variable marks the field as either 1 or 0
buyerType=pd.get_dummies(heavyCat)
buyerType[:3]
```

```
[11]:      regular  heavy
0         1      0
1         0      1
2         1      0
```

```
[12]: # create new columns
xyzcustnew['typeReg']=buyerType['regular']
xyzcustnew['typeHeavy']=buyerType['heavy']
```

```
[13]: xyzcustnew.columns
```

```
[13]: Index(['index', 'ACCTNO', 'ZIP', 'ZIP4', 'LTD_SALES', 'LTD_TRANSACTIONS',
          'YTD_SALES_2009', 'YTD_TRANSACTIONS_2009', 'CHANNEL_ACQUISITION',
          'BUYER_STATUS', 'ZIP9_SUPERCODE', 'heavyCat', 'typeReg', 'typeHeavy'],
          dtype='object')
```

```
[14]: # look at new variables
xyzcustnew.head()
```

```
[14]:      index  ACCTNO  ZIP  ZIP4  LTD_SALES  LTD_TRANSACTIONS  YTD_SALES_2009  \
0         0  WDQQLDQL  60084  5016         90.0              1          0.0
1         1  WQWAYHYLA  60091  1750        4227.0             9        1263.0
2         2  GSHAPLHAW  60067   900         420.0             3         129.0
3         3  PGGYDYWAD  60068  3838        6552.0             6          0.0
4         4  LWPSGPLLS  60090  3932         189.0             3          72.0
```

```
      YTD_TRANSACTIONS_2009  CHANNEL_ACQUISITION  BUYER_STATUS  ZIP9_SUPERCODE  \
0                        0                IB      INACTIVE      600845016
1                        3                RT      ACTIVE      600911750
2                        1                RT      ACTIVE      600670900
3                        0                RT      INACTIVE      600683838
4                        1                RT      ACTIVE      600903932
```

```
      heavyCat  typeReg  typeHeavy
0  regular      1         0
1   heavy      0         1
2  regular      1         0
3  regular      1         0
4  regular      1         0
```

```
[15]: # for this exercises we need to create trCountsChrono object
```

```
xyztrans=pd.read_sql('xyztrans', engine)

trandate=xyztrans.TRANDATE      # should be a Series
```

```

daystr=trandate.str[0:2]          # two digit date numbers slice

mostr=trandate.str[2:5]          # the three letter month abbreviations

yearstr=trandate.str[5:]         # four digit years

print(daystr[0],mostr[0],yearstr[0],xyztrans.TRANDATE[0])

```

09 JUN 2009 09JUN2009

```

[16]: #create a dictionary for the months
monums={'JAN':'1', 'FEB':'2', 'MAR':'3', 'APR':'4', 'MAY':'5', 'JUN':'6',
        'JUL':'7', 'AUG':'8', 'SEP':'9', 'OCT':'10', 'NOV':'11', 'DEC':'12'}
#month
monos=mostr.map(monums)          # do a dict lookup for each value of mostr

transtr=yearstr+'-'+monos+'-'+daystr
print(transtr[0])

```

2009-6-09

transtr should be a Series. Now let's convert the string values in transtr into datetime values:

```

[17]: # convert to datetime values
trDateTime=pd.to_datetime(transtr)

```

```

[18]: trCounts=trDateTime.value_counts()
trCounts

```

```

[18]: 2009-12-19    877
      2009-12-21    836
      2009-12-12    782
      2009-12-23    765
      2009-12-20    744
      2009-12-22    717
      2009-12-18    708
      2009-12-14    615
      2009-12-15    599
      2009-12-16    571
      2009-12-11    568
      2009-11-21    561
      2009-12-13    542
      2009-11-22    507
      2009-12-10    504
      2009-12-04    488
      2009-11-25    451
      2009-12-24    425
      2009-11-23    421
      2009-11-27    419

```

2009-11-24	412
2009-04-10	404
2009-11-28	402
2009-11-14	402
2009-12-09	401
2009-05-09	398
2009-12-08	397
2009-11-07	394
2009-12-07	372
2009-01-17	372
...	
2009-02-26	47
2009-04-30	47
2009-02-21	39
2009-03-01	36
2009-03-10	33
2009-06-14	29
2009-11-26	26
2009-03-16	24
2009-07-14	24
2009-06-15	24
2009-08-30	23
2009-11-04	23
2009-02-25	22
2009-02-02	22
2009-10-17	21
2009-03-08	20
2009-03-27	19
2009-02-03	19
2009-02-01	18
2009-03-11	17
2009-02-07	16
2009-07-02	15
2009-06-13	14
2009-07-16	14
2009-12-25	11
2009-08-15	11
2009-04-12	10
2009-10-13	10
2009-08-21	6
2009-03-15	5

Name: TRANDATE, Length: 365, dtype: int64

The order of the counts in `trDateTime` is not chronological, so let's reorder them so that they go from earliest to most recent date.

```
[19]: newIndex=pd.date_range(trCounts.index.min(),trCounts.index.max())
```

```
trCountsChrono=trCounts.reindex(index=newIndex)
```

```
[20]: print(trCountsChrono.head())
```

```
2009-01-01    176
2009-01-02    305
2009-01-03    365
2009-01-04    231
2009-01-05    144
Freq: D, Name: TRANDATE, dtype: int64
```

One of the very handy things you can do with pandas DataFrames and Series is that you can create what are called hierarchical indexes. These are multi-level indexes (they are in fact called MultiIndexes). They make it easier to select, modify, group, and reshape data in a wide variety of ways. They make it possible to work with high dimensional data in data structures that are in just one or two dimensions.

Let's change trCountsChrono a bit to produce a first simple example of a Series with a hierarchical index. First, let's put the Series into a DataFrame and then rename the columns:

```
[21]: # initialize a dataframe
trDF=DataFrame()
```

```
[22]: trDF
```

```
[22]: Empty DataFrame
Columns: []
Index: []
```

```
[23]: # load dataframe with 2 columns
trDF['date'] = trCountsChrono.index
trDF['transactions'] = trCountsChrono.values
trDF.columns
```

```
[23]: Index(['date', 'transactions'], dtype='object')
```

```
[24]: trDF.head()
```

```
[24]:
```

	date	transactions
0	2009-01-01	176
1	2009-01-02	305
2	2009-01-03	365
3	2009-01-04	231
4	2009-01-05	144

```
[25]: trDF.dtypes
```

```
[25]: date                datetime64[ns]
transactions          int64
dtype: object
```

Note that the data types of the columns have not changed. Try trDF.dtypes.

Now, let's create a new column that indicates whether the number of daily transactions are heavy or light depending on whether they are equal to or greater than the median number of trans-

actions, or less than the median number. There are more succinct ways to do this, but this is transparent, if not efficient:

```
[26]: trMed=trDF.transactions.median()           # here's the median
      trMed
```

```
[26]: 136.0
```

```
[27]: # if the value is greater than or equal to the median, then heavy
      heavyLight = lambda x : x >= trMed and 'heavy' or 'light' # an example anon_
      ↪function
```

```
[28]: # use map to call lambda
      trDF['vol']=trDF.transactions.map(heavyLight)           # 'vol' is the heavy/light_
      ↪column
      trDF.head(10)
```

```
[28]:      date  transactions  vol
0 2009-01-01           176  heavy
1 2009-01-02           305  heavy
2 2009-01-03           365  heavy
3 2009-01-04           231  heavy
4 2009-01-05           144  heavy
5 2009-01-06           188  heavy
6 2009-01-07           166  heavy
7 2009-01-08            52  light
8 2009-01-09           194  heavy
9 2009-01-10           166  heavy
```

Note that this lambda would stumble if trMed wasn't known at the time lambda was called by the map method.

Anyway, next we're going to create, monum, a variable indicating the month of the calendar year that each day falls into:

```
[29]: trDF['monum'] = trDF.date.dt.month           # .dt is the datetime accessor
```

Next, we're going to collapse the daily transaction counts into monthly counts. When we do this we'll keep the heavy versus light daily volume distinction. First we're going to drop the 'date' column because we no longer need it. To be safe we'll copy the result to a new DataFrame just in case something goes wrong:

```
[30]: # making a copy of trDF while also dropping date as a column
      # axis=1 means here a column is selected to drop
      trDFnd=trDF.drop('date',axis=1)
      trDFnd.head()
```

```
[30]:      transactions  vol  monum
0           176  heavy      1
1           305  heavy      1
2           365  heavy      1
3           231  heavy      1
4           144  heavy      1
```

Now using this DataFrame's groupby() method, sum up the transactions within month by heavy

volume days and light volume days:

```
[31]: trDFgrouped = trDFnd.groupby(['monum', 'vol']).sum()
trDFgrouped.head()
```

```
[31]:                transactions
monum vol
1    heavy          5255
     light           572
2    heavy           761
     light          1625
3    heavy          1130
```

```
[32]: # check out the indexes
trDFgrouped.index.levels
```

```
[32]: FrozenList([[1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12], ['heavy', 'light']])
```

Now if you look at this DataFrame you'll see that it has two levels of indexing, monum, and within the levels of monum, vol. If you enter `trDFgrouped.index` you'll get back a MultiIndex object. Also, try `trDFgrouped.index.levels` to see what you get.

pandas has pretty seamlessly created this index for you, but you can construct MultiIndexes manually by combining equal length arrays (using `MultiIndex.from_arrays`) of index levels, or by using tuples (with `MultiIndex.from_tuples`). In both cases all combinations of the index levels are included.

Note that if you look at `trDFgrouped` you may see here and there that for a particular month, the number of heavy day transactions is less than the number of light day transactions. How do you think that could happen?

You can use MultiIndexes to select and subset DataFrames and Series in many of the same ways you can use simple indexes. For example, to get the heavy days transaction count data for November, you can do:

```
[33]: trDFgrouped.loc[11, 'heavy']
```

```
[33]: transactions      8402
Name: (11, heavy), dtype: int64
```

```
[34]: # first six months of data
trDFgrouped.loc[list(range(1,7))]
```

```
[34]:                transactions
monum vol
1    heavy          5255
     light           572
2    heavy           761
     light          1625
3    heavy          1130
     light          1664
4    heavy          2327
     light          1727
5    heavy          2172
     light          2076
6    heavy          2878
```

light	1495
-------	------

```
[35]: # or the first 6 rows of data:
trDFgrouped.iloc[0:6]                # .iloc here, but .loc above.
```

```
[35]: transactions
```

monum	vol
1	heavy
	5255
	light
	572
2	heavy
	761
	light
	1625
3	heavy
	1130
	light
	1664

The data starting from the March heavy day counts to the July light counts:

```
[36]: trDFgrouped[(3, 'light'):(7, 'heavy')]
```

```
[36]: transactions
```

monum	vol
3	light
	1664
4	heavy
	2327
	light
	1727
5	heavy
	2172
	light
	2076
6	heavy
	2878
	light
	1495
7	heavy
	4440

The above uses a range defined by a slice of tuples. So does:

```
[37]: trDFgrouped[(3, 'light'):6]
```

```
[37]: transactions
```

monum	vol
3	light
	1664
4	heavy
	2327
	light
	1727
5	heavy
	2172
	light
	2076
6	heavy
	2878
	light
	1495

Try selecting some data and slicing a few times yourself. It takes a little practice to get the hang of getting what you want.

There are many other ways to slice using MultiIndexes. One other you might find interesting is the cross-section method `.xs`. Here's an example that picks out data for the light days:

```
[38]: trDFgrouped.xs('light', level='vol')
```

```
[38]: transactions
```

monum
1
572

2	1625
3	1664
4	1727
5	2076
6	1495
7	564
8	1938
9	1942
10	2241
11	49
12	257

As you probably know, DataFrames have a transpose method, .T:

```
[39]: trDFgrouped.xs('light',level='vol').T           # the transpose of the
      ↪above
```

```
[39]: monum      1      2      3      4      5      6      7      8      9     10    11  \
      transactions  572  1625  1664  1727  2076  1495  564  1938  1942  2241  49
```

```
monum      12
transactions 257
```

Did you get a table of transactions with cells labeled by monum across the top?

You can also pivot DataFrames in various ways. Let's make some data to create a DataFrame we can pivot. We'll put the monum and vol indexes from trDFgrouped into our new DataFrame as columns, and then we'll add transactions as a third column.

```
[40]: mo=trDFgrouped.index.get_level_values(0)       # the month numbers
```

```
[41]: volType=trDFgrouped.index.get_level_values(1) # vol
```

```
[42]: trDFpiv=DataFrame({'month':mo,'vol': volType, 'transactions':trDFgrouped.
      ↪transactions})           # data as a dict
```

```
[43]: trDFpiv
```

```
[43]:
```

		month	transactions	vol
monum	vol			
1	heavy	1	5255	heavy
	light	1	572	light
2	heavy	2	761	heavy
	light	2	1625	light
3	heavy	3	1130	heavy
	light	3	1664	light
4	heavy	4	2327	heavy
	light	4	1727	light
5	heavy	5	2172	heavy
	light	5	2076	light
6	heavy	6	2878	heavy
	light	6	1495	light
7	heavy	7	4440	heavy

	light	7	564	light
8	heavy	8	1682	heavy
	light	8	1938	light
9	heavy	9	1921	heavy
	light	9	1942	light
10	heavy	10	2109	heavy
	light	10	2241	light
11	heavy	11	8402	heavy
	light	11	49	light
12	heavy	12	13168	heavy
	light	12	257	light

Now, let's pivot trDFpiv. Let's make a new DataFrame with month as the index, vol the columns, and the transaction counts as the values:

```
[44]: trDFpived = trDFpiv.pivot(index='month', columns='vol', values='transactions')
trDFpived
```

```
[44]: vol    heavy    light
month
1      5255     572
2       761    1625
3      1130    1664
4      2327    1727
5      2172    2076
6      2878    1495
7      4440     564
8      1682    1938
9      1921    1942
10     2109    2241
11     8402      49
12    13168     257
```

How does trDFpived look to you?

If trDFpiv had more than one column for values not used as a column or an index, hierarchical columns would be created to reflect them. For example, let's add an additional column to trDFpiv:

```
[45]: trDFpiv['randy'] = np.random.randn(len(trDFpiv))
trDFpiv.head()
```

```
[45]:
```

		month	transactions	vol	randy
monum	vol				
1	heavy	1	5255	heavy	2.072689
	light	1	572	light	1.551954
2	heavy	2	761	heavy	1.611733
	light	2	1625	light	-0.551663
3	heavy	3	1130	heavy	0.224400

Now pivot trDFpiv like:

```
[46]: trDFpived2 = trDFpiv.pivot(index='month', columns='vol')
trDFpived2.head()
```

```
[46]: transactions      randy
      vol      heavy light      heavy      light
      month
      1          5255   572  2.072689  1.551954
      2           761  1625  1.611733 -0.551663
      3          1130  1664  0.224400  1.429708
      4          2327  1727 -0.063512  1.104748
      5          2172  2076  1.336558 -1.077305
```

How does trDFpived2 look?

OK, let's drop randy from trDFpiv and try some other things.

Feeling lucky? Then do `trDFpiv.drop('randy',axis=1,inplace=True)`.

You can also stack and unstack DataFrames. These methods come in handy when you need to shape some data in a particular way to be input to an algorithm. Let's aggregate some of the xyzcustnew data (see above) to get a DataFrame we can stack and unstack:

```
[47]: # remember we read in xyzcustnew from xyz.db
      xyzdata = xyzcustnew[['BUYER_STATUS', 'heavyCat', 'CHANNEL_ACQUISITION']]
      xyzdata.head()
```

```
[47]: BUYER_STATUS heavyCat CHANNEL_ACQUISITION
      0      INACTIVE  regular                IB
      1        ACTIVE   heavy                RT
      2        ACTIVE  regular                RT
      3      INACTIVE  regular                RT
      4        ACTIVE  regular                RT
```

Use xyzdata because it's just easier. It has just the three columns we're now going to work with.

```
[48]: xyzgrouped = xyzdata.groupby(['BUYER_STATUS', 'heavyCat', 'CHANNEL_ACQUISITION'])
```

```
[49]: xyzCountData = xyzgrouped.size()           # a MultiIndexed Series of counts
      xyzCountData
```

```
[49]: BUYER_STATUS  heavyCat  CHANNEL_ACQUISITION
      ACTIVE      regular      CB                443
      ACTIVE      regular      IB               1112
      ACTIVE      regular      RT               7393
      ACTIVE      heavy      CB                356
      ACTIVE      heavy      IB                703
      ACTIVE      heavy      RT               3325
      INACTIVE      regular      CB                691
      INACTIVE      regular      IB               1249
      INACTIVE      regular      RT               7056
      LAPSED      regular      CB                372
      LAPSED      regular      IB               1111
      LAPSED      regular      RT               6368

      dtype: int64
```

```
[50]: print(xyzCountData.unstack())
```

```
CHANNEL_ACQUISITION      CB      IB      RT
```

BUYER_STATUS	heavyCat			
ACTIVE	regular	443	1112	7393
	heavy	356	703	3325
INACTIVE	regular	691	1249	7056
LAPSED	regular	372	1111	6368

xyzCountData is a Series with a MultiIndex, and so it can be unstacked, changing it from tall and narrow to short and wide. Note that by default, only the lowest level of the MultiIndex is used for unstacking. Do you know why there are no heavy buyers in the INACTIVE or LAPSED categories?

Let's restack this into a different version of xyzCountData:

```
[51]: unStackxyz = xyzCountData.unstack()           # what we had just above
      unStackxyz
```

```
[51]: CHANNEL_ACQUISITION    CB    IB    RT
      BUYER_STATUS heavyCat
      ACTIVE      regular    443    1112    7393
              heavy      356     703    3325
      INACTIVE      regular    691    1249    7056
      LAPSED        regular    372    1111    6368
```

```
[52]: unStackxyz.T.stack()           # .T is the transpose
```

```
[52]: BUYER_STATUS              ACTIVE  INACTIVE  LAPSED
      CHANNEL_ACQUISITION heavyCat
      CB                  regular      443      691.0    372.0
              heavy      356         NaN      NaN
      IB                  regular     1112     1249.0   1111.0
              heavy      703         NaN      NaN
      RT                  regular     7393     7056.0   6368.0
              heavy     3325         NaN      NaN
```

Note how in the above, combinations of the levels of the three variables that do not actually occur in the data are given an NaN, a missing value. NaN means not a number. The cells are stacked using levels of BUYER_STATUS within level of CHANNEL_ACQUISITION.

Try doing unStackxyz.T.stack(1) to get stacking by heavyCat instead of by BUYER_STATUS. Here again, cells do not have observations are given a NaN.

The unstack method can return a stacked object as it was when it was stacked, but it can also return it in a different unstacked form. For example, see what this does:

```
[53]: unStackxyz.T.stack(0).unstack(1)
```

```
[53]: heavyCat      regular      heavy
      BUYER_STATUS    ACTIVE  INACTIVE  LAPSED    ACTIVE  INACTIVE  LAPSED
      CHANNEL_ACQUISITION
      CB              443      691      372    356.0      NaN      NaN
      IB              1112     1249     1111    703.0      NaN      NaN
      RT              7393     7056     6368   3325.0      NaN      NaN
```

You can stack or unstack on multiple levels at one time. See what this does for you:

```
[54]: unStackxyz.T.stack(level=['heavyCat', 'BUYER_STATUS'])
```

```
[54]: CHANNEL_ACQUISITION heavyCat BUYER_STATUS
CB          regular    ACTIVE      443.0
          regular    INACTIVE     691.0
          regular    LAPSED       372.0
          heavy     ACTIVE      356.0
IB          regular    ACTIVE     1112.0
          regular    INACTIVE     1249.0
          regular    LAPSED     1111.0
          heavy     ACTIVE       703.0
RT          regular    ACTIVE     7393.0
          regular    INACTIVE     7056.0
          regular    LAPSED     6368.0
          heavy     ACTIVE     3325.0

dtype: float64
```

and compare to:

```
[55]: unStackxyz.T.stack(level=['BUYER_STATUS', 'heavyCat'])
```

```
[55]: CHANNEL_ACQUISITION BUYER_STATUS heavyCat
CB          ACTIVE      regular      443.0
          ACTIVE      heavy        356.0
          INACTIVE     regular      691.0
          LAPSED       regular      372.0
IB          ACTIVE     regular     1112.0
          heavy        heavy        703.0
          INACTIVE     regular     1249.0
          LAPSED       regular     1111.0
RT          ACTIVE     regular     7393.0
          heavy        heavy     3325.0
          INACTIVE     regular     7056.0
          LAPSED       regular     6368.0

dtype: float64
```

The pandas melt() method provides some similar functionality. You can use it to turn a short and wide DataFrame into a taller, narrower one by identifying columns that contain values to be used as record identifiers. Let's go back to the xyzcustnew data and select a few columns from it to do some melting on:

```
[56]: xyzcust = xyzcustnew[['BUYER_STATUS', 'heavyCat', 'LTD_SALES']].copy()
```

Now, let's melt xyzcust so that BUYER_STATUS and heavyCat become identifiers :

xyzcustm will look something like:

```
[58]: print(xyzcustm)
```

```
   BUYER_STATUS heavyCat LTD_SALES  value
0      INACTIVE  regular  LTD_SALES    90.0
1      ACTIVE    heavy   LTD_SALES  4227.0
```

2	ACTIVE	regular	LTD_SALES	420.0
3	INACTIVE	regular	LTD_SALES	6552.0
4	ACTIVE	regular	LTD_SALES	189.0
5	ACTIVE	regular	LTD_SALES	4278.0
6	ACTIVE	heavy	LTD_SALES	1869.0
7	ACTIVE	regular	LTD_SALES	33.0
8	INACTIVE	regular	LTD_SALES	735.0
9	INACTIVE	regular	LTD_SALES	468.0
10	ACTIVE	regular	LTD_SALES	804.0
11	LAPSED	regular	LTD_SALES	219.0
12	ACTIVE	heavy	LTD_SALES	3240.0
13	INACTIVE	regular	LTD_SALES	180.0
14	ACTIVE	regular	LTD_SALES	423.0
15	INACTIVE	regular	LTD_SALES	306.0
16	LAPSED	regular	LTD_SALES	1002.0
17	ACTIVE	regular	LTD_SALES	1155.0
18	ACTIVE	regular	LTD_SALES	612.0
19	ACTIVE	regular	LTD_SALES	633.0
20	INACTIVE	regular	LTD_SALES	114.0
21	ACTIVE	regular	LTD_SALES	294.0
22	INACTIVE	regular	LTD_SALES	849.0
23	INACTIVE	regular	LTD_SALES	72.0
24	ACTIVE	heavy	LTD_SALES	3411.0
25	ACTIVE	heavy	LTD_SALES	1023.0
26	LAPSED	regular	LTD_SALES	873.0
27	ACTIVE	heavy	LTD_SALES	2778.0
28	ACTIVE	heavy	LTD_SALES	2676.0
29	LAPSED	regular	LTD_SALES	528.0
...
30149	ACTIVE	regular	LTD_SALES	861.0
30150	ACTIVE	regular	LTD_SALES	837.0
30151	ACTIVE	regular	LTD_SALES	2478.0
30152	ACTIVE	regular	LTD_SALES	84.0
30153	ACTIVE	heavy	LTD_SALES	2877.0
30154	INACTIVE	regular	LTD_SALES	1611.0
30155	LAPSED	regular	LTD_SALES	1860.0
30156	LAPSED	regular	LTD_SALES	48.0
30157	ACTIVE	regular	LTD_SALES	195.0
30158	LAPSED	regular	LTD_SALES	60.0
30159	INACTIVE	regular	LTD_SALES	252.0
30160	LAPSED	regular	LTD_SALES	594.0
30161	LAPSED	regular	LTD_SALES	1272.0
30162	ACTIVE	heavy	LTD_SALES	2184.0
30163	ACTIVE	regular	LTD_SALES	759.0
30164	INACTIVE	regular	LTD_SALES	756.0
30165	ACTIVE	regular	LTD_SALES	1365.0
30166	ACTIVE	heavy	LTD_SALES	2490.0
30167	ACTIVE	heavy	LTD_SALES	438.0

30168	INACTIVE	regular	LTD_SALES	549.0
30169	ACTIVE	regular	LTD_SALES	150.0
30170	ACTIVE	regular	LTD_SALES	93.0
30171	INACTIVE	regular	LTD_SALES	834.0
30172	INACTIVE	regular	LTD_SALES	147.0
30173	LAPSED	regular	LTD_SALES	816.0
30174	ACTIVE	regular	LTD_SALES	2736.0
30175	ACTIVE	regular	LTD_SALES	2412.0
30176	INACTIVE	regular	LTD_SALES	429.0
30177	INACTIVE	regular	LTD_SALES	651.0
30178	ACTIVE	heavy	LTD_SALES	4527.0

[30179 rows x 4 columns]

You'll probably realize that the leftmost column is a simple numerical index that this pandas method created. There's a pandas method called `wide_table` that works similarly, but can be a little easier to use. Give it a try using `xyzcust` or the `DataFrame` of your choice.

So at this point we've pivoted, grouped, and reshaped. The pivoting example we did was pretty simple. pandas also provides a method called `pivot_table` that provides considerable flexibility in terms of how data can be reorganized and summarized. Let's consider the example below.

```
[59]: pd.
      →pivot_table(xyzcustnew, values='YTD_SALES_2009', index=['BUYER_STATUS', 'heavyCat'], columns=['CHANNEL_ACQUISITION', 'CB', 'IB', 'RT'])
```

```
[59]: CHANNEL_ACQUISITION      CB      IB      RT
BUYER_STATUS heavyCat
ACTIVE      regular      205.334086      191.047662      167.993913
           heavy      2397.606742      1251.559033      1158.506165
INACTIVE    regular      0.000000      0.000000      0.000000
           heavy      NaN      NaN      NaN
LAPSED      regular      0.000000      0.000000      0.000000
           heavy      NaN      NaN      NaN
```

Do you see some rows in the result that only have zeros? Why are they there?

Or, try doing:

```
[60]: pd.
      →pivot_table(xyzcustnew, values='YTD_SALES_2009', index=['BUYER_STATUS'], columns=['heavyCat', 'CHANNEL_ACQUISITION', 'CB', 'IB', 'RT'])
```

```
[60]: heavyCat      regular      heavy \
CHANNEL_ACQUISITION      CB      IB      RT      CB
BUYER_STATUS
ACTIVE      205.334086      191.047662      167.993913      2397.606742
INACTIVE      0.000000      0.000000      0.000000      NaN
LAPSED      0.000000      0.000000      0.000000      NaN

heavyCat
CHANNEL_ACQUISITION      IB      RT
BUYER_STATUS
ACTIVE      1251.559033      1158.506165
INACTIVE      NaN      NaN
```

NaN

`pivot_table` defaults to taking the mean (using `np.mean`) of the groups it defines. If you want some other aggregation instead

np.sum :

```
→ index= ['BUYER_STATUS'], columns= ['heavyCat', 'CHANNEL_ACQUISITION'], aggfunc=np.  
→ sum)
```

heavyCat	
CHANNEL_ACQUISITION	RT
BUYER_STATUS	
ACTIVE	3852033.0
INACTIVE	NaN
LAPSED	NaN

You can also add margins to pivot tables by using the margins = True option. For example, to get row and column totals :

heavyCat		All
CHANNEL_ACQUISITION	RT	
BUYER_STATUS		
ACTIVE	3852033.0	7130814.0
INACTIVE	NaN	0.0
LAPSED	NaN	0.0
All	3852033.0	7130814.0

Should give you the same table as above but with row and column totals added.

It has probably dawned on you that you can manipulate data objects in many different ways to group them and to apply descriptive statistics to them. Let's group xyz customers using `BUYER5TATUSandheavyCat` :

```
[63]: xyzGrouper=xyzcustnew.groupby(['BUYER STATUS', 'heavyCat'])
```

groupby can apply conventional as well as custom functions to aggregated data. For example:

```
[64]: xyzGrouper.agg({'YTD_SALES_2009': [np.mean, np.std], 'LTD_SALES': [np.mean, np.
    ↳std]})
```

```
[64]:
```

		YTD_SALES_2009		LTD_SALES	
		mean	std	mean	std
BUYER_STATUS	heavyCat				
ACTIVE	regular	172.707532	107.584023	1001.845105	1466.075631
	heavy	1274.048130	5434.616517	4096.179745	34210.646330
INACTIVE	regular	0.000000	0.000000	568.014784	850.966479
LAPSED	regular	0.000000	0.000000	841.467329	1374.447756

calculates the mean and standard deviation of YTD_SALES₂₀₀₉ and LTD_SALES for each of the groups defined in xyzGr

Try using a version of this command to get statistics for the columns YTD_TRANS_{ACTIONS}₂₀₀₉ and LTD_TRANS_{ACTIONS}. These are both count variables. What descriptive statistics do you

Note that you can apply custom functions to data aggregates. Suppose we wanted to compute the coefficient of variation, CV, for data. The CV is a standardized measure of dispersion, and is the ratio of the standard deviation to the mean. It's estimated by the ratio of the estimates of these two statistics. We could write our own function to do this:

```
[65]: def coefV(x): # a baby CV function that accepts a
    ↳sequence
    return np.std(x)/np.mean(x)
```

This will work assuming that the mean and std numpy methods are available in this function's namespace, of course. Note that our baby function doesn't do anything smart regarding missing values and other inconveniences, but it's good enough to demonstrate what we want, here. What do you think it means if what it produces is negative? How could that happen?

We can apply this function to selected groups. Here we apply it to customers grouped by BUYER_STATUS. Let's first get a simpler DataFrame to fiddle with:

```
[66]: buyerStats=xyzcustnew[['BUYER_STATUS', 'LTD_SALES', 'LTD_TRANSACTIONS']]
buyerGrouper=buyerStats.groupby(['BUYER_STATUS'])
buyerGrouper.agg(coefV)
```

```
[66]:
```

	LTD_SALES	LTD_TRANSACTIONS
BUYER_STATUS		
ACTIVE	9.758480	1.153501
INACTIVE	1.498058	0.784441
LAPSED	1.633290	0.987139

Did you get a table of CV's?

We could combine our own function or functions with existing functions and apply them on a group by group basis. Let's play with a function that returns 5th and 95th percentiles of some data:

```
[67]: def ptils(x):
    p5 = np.percentile(x,5)
    p95 = np.percentile(x,95)
    return p5, p95
```

There's our toy function. coefV, it may break with bad data. (So, watch out.)

What kind of object does ptils return?

Now, applying np.mean and ptils:

```
[68]: buyerGrouper.agg([np.mean, ptils])
```

```
[68]:
```

	LTD_SALES		LTD_TRANSACTIONS	\
	mean		ptils	mean
BUYER_STATUS				
ACTIVE	2019.364086	(81.0, 6544.349999999997)		6.935794
INACTIVE	568.014784	(60.0, 1776.0)		2.263895
LAPSED	841.467329	(63.0, 2904.0)		3.498280


```
ptils
```

BUYER_STATUS	
ACTIVE	(1.0, 20.0)
INACTIVE	(1.0, 6.0)
LAPSED	(1.0, 9.0)

What kind of object is the above command printing out for you?

You can select particular results from this, of course, e.g.:

```
[69]: buyerGrouper.agg([np.mean, ptils]).loc['ACTIVE', 'LTD_SALES']
```

```
[69]: mean                2019.36
ptils    (81.0, 6544.349999999997)
Name: ACTIVE, dtype: object
```

As a quick little exercise to do on you own, write a tiny function that calculates the interquartile range, or IQR, for data, and then apply it to the above data. The IQR is the difference between the 75th and the 25th percentile values.

Well, that wraps it up for this, and last, Python Practice. No surprisingly, there's a lot more to data management using Python and packages like Pandas, and there's something new all the time.

If you're an R user, and you use it on Linux or OS X, you'll want to check out the package rpy2, which provides some capability for transferring data between R and Python. It's under development, and the plan is that it will eventually allow doing things like calling R functions from within Python. It is apparently pretty tough to install and use from in Windows at the present time.

3 Requirements :

1. Get the trDFgrouped data starting from the May heavy day counts to the August heavy counts
2. Group xyz customers using BUYER_STATUS, heavyCat, and ZIP, and apply np.sum function on the aggregated data for YTD_SALES_2009 and LTD_SALES columns

```
[70]: # Write your python code that meets the above requirements in this cell
# Question 1: Get the trDFgrouped data starting from the May heavy day counts
→to the August heavy counts
trDFgrouped[(5, 'heavy'):(8, 'heavy')]
```

```
[70]: transactions
monum vol
5     heavy      2172
      light      2076
6     heavy      2878
      light      1495
7     heavy      4440
      light       564
8     heavy      1682
```

```
[71]: # Question 2: Group xyz customers using BUYER_STATUS, heavyCat, and ZIP,
# and apply np.sum function on the aggregated data for YTD_SALES_2009 and
      ↳ LTD_SALES columns
xyzGrouper = xyzcustnew.groupby(['BUYER_STATUS', 'heavyCat', 'ZIP'])
xyzGrouper.agg({'YTD_SALES_2009': [np.sum], 'LTD_SALES': [np.sum]})
```

```
[71]: YTD_SALES_2009 LTD_SALES
      sum sum
BUYER_STATUS heavyCat ZIP
ACTIVE      regular 60056 68913.0 332196.0
              60060 68520.0 339567.0
              60061 68328.0 400569.0
              60062 141237.0 762387.0
              60064 2169.0 9129.0
              60065 1002.0 2784.0
              60067 156429.0 922680.0
              60068 140133.0 802815.0
              60069 43623.0 280686.0
              60070 24051.0 134265.0
              60071 4311.0 20112.0
              60072 2037.0 14583.0
              60073 29877.0 143901.0
              60074 72999.0 349026.0
              60076 53040.0 252438.0
              60077 39546.0 183588.0
              60078 1878.0 7410.0
              60081 16446.0 76662.0
              60083 14445.0 81954.0
              60084 39834.0 243837.0
              60085 18714.0 88857.0
              60087 13749.0 59997.0
              60088 1053.0 2538.0
              60089 100038.0 481086.0
              60090 32934.0 153108.0
              60091 178533.0 1127982.0
              60093 169671.0 1449606.0
              60094 357.0 543.0
              60096 5544.0 34929.0
```

		60097	5805.0	29565.0
...		
LAPSED	regular	60064	0.0	3537.0
		60065	0.0	7359.0
		60067	0.0	682167.0
		60068	0.0	571056.0
		60069	0.0	134685.0
		60070	0.0	75333.0
		60071	0.0	11232.0
		60072	0.0	2463.0
		60073	0.0	100932.0
		60074	0.0	245877.0
		60076	0.0	207912.0
		60077	0.0	135801.0
		60078	0.0	4173.0
		60079	0.0	2928.0
		60081	0.0	50397.0
		60082	0.0	225.0
		60083	0.0	71463.0
		60084	0.0	157020.0
		60085	0.0	60144.0
		60087	0.0	45030.0
		60088	0.0	3354.0
		60089	0.0	407976.0
		60090	0.0	137544.0
		60091	0.0	820053.0
		60093	0.0	955428.0
		60095	0.0	300.0
		60096	0.0	17559.0
		60097	0.0	30564.0
		60098	0.0	149418.0
		60192	0.0	4548.0

[132 rows x 2 columns]